

# CMSC 3540I: The Interplay of Learning and Game Theory (Autumn 2022)

## Introduction

---



Instructor: Haifeng Xu

# Outline

- Course Overview
- Administrivia
- An Example

# Single-Agent Decision Making

- A decision maker picks an action  $x \in X$ , resulting in utility  $f(x)$
- Typically an optimization problem:

$$\begin{array}{ll}\text{minimize (or maximize)} & f(x) \\ \text{subject to} & x \in X\end{array}$$

- $x$ : decision variable
- $f(x)$ : objective function
- $X$ : feasible set/region
- Optimal solution, optimal value

- Example 1: minimize  $x^2$ , s.t.  $x \in [-1,1]$

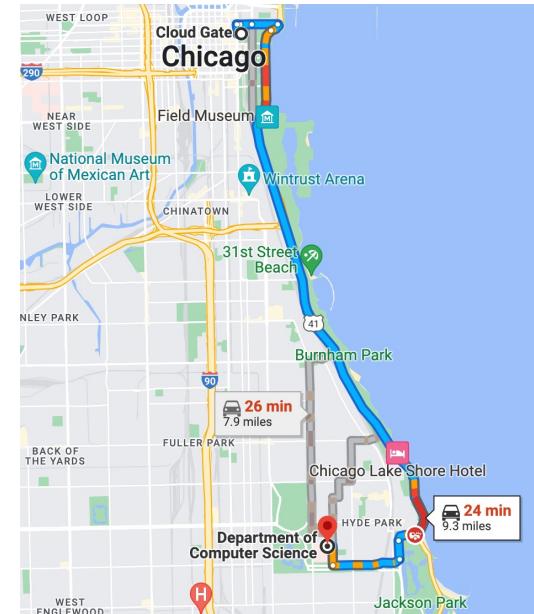
# Single-Agent Decision Making

- A decision maker picks an action  $x \in X$ , resulting in utility  $f(x)$
- Typically an optimization problem:

$$\begin{array}{ll} \text{minimize (or maximize)} & f(x) \\ \text{subject to} & x \in X \end{array}$$

- $x$ : decision variable
- $f(x)$ : objective function
- $X$ : feasible set/region
- Optimal solution, optimal value

- Example 1: minimize  $x^2$ , s.t.  $x \in [-1,1]$
- Example 2: pick a road to school



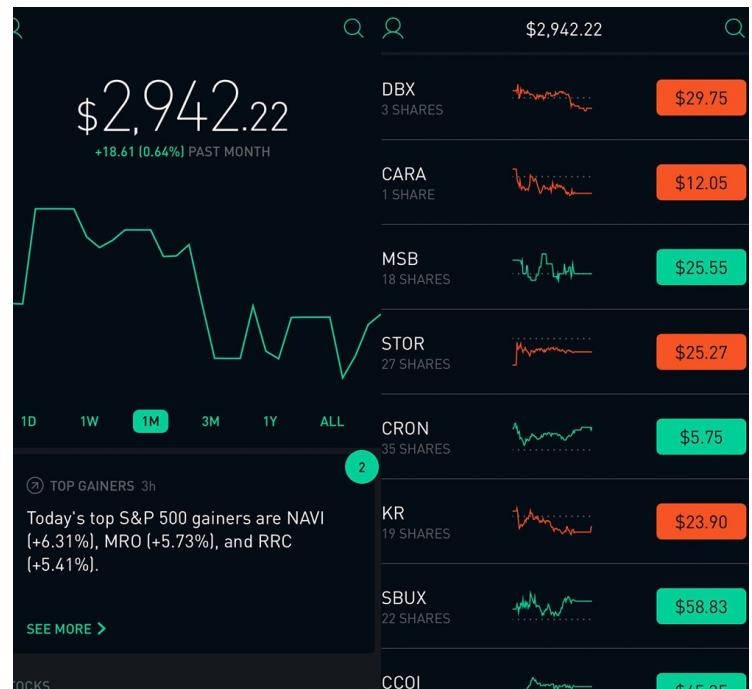
# Single-Agent Decision Making

- A decision maker picks an action  $x \in X$ , resulting in utility  $f(x)$
- Typically an **optimization problem**:

$$\begin{array}{ll} \text{minimize (or maximize)} & f(x) \\ \text{subject to} & x \in X \end{array}$$

- $x$ : decision variable
- $f(x)$ : objective function
- $X$ : feasible set/region
- Optimal solution, optimal value

- Example 1: minimize  $x^2$ , s.t.  $x \in [-1,1]$
- Example 2: pick a road to school
- Example 3: invest a subset of stocks



# Multi-Agent Decision Making

- Usually, your payoffs affected not only by your actions, but also others'
- Agent  $i$ 's utility  $f_i(x_i, x_{-i})$  depends on his own action  $x_i$ , as well as other agents' actions  $x_{-i}$
- Is this still an optimization problem? Should each agent  $i$  just pick  $x_i \in X_i$  to minimize  $f_i(x_i, x_{-i})$ ?
  - $x_{-i}$  is not under  $i$ 's control
  - Think of rock-paper-scissor game
- Examples: stock investment, routing, sales, even taking courses...

# Example I: Prisoner's Dilemma

- Two members A,B of a criminal gang are arrested
- They are questioned in two separate rooms
  - ❖ No communications between them



	B	B stays silent	B betrays
A			
A stays silent	-1	0	
A betrays	-3	-2	

Q: How should each prisoner act?

- Betray is always the best action

# Example I: Prisoner's Dilemma

- Two members A,B of a criminal gang are arrested
- They are questioned in two separate rooms
  - ❖ No communications between them



	B	B stays silent	B betray
A	B stays silent	-1	0
A betray	-1	-3	-2

Q: How should each prisoner act?

- Betray is always the best action

# Example I: Prisoner's Dilemma

- Two members A,B of a criminal gang are arrested
- They are questioned in two separate rooms
  - ❖ No communications between them

	B	B stays silent	B betrays
A stays silent	-1	-3	0
A betrays	0	-3	-2

equilibrium

Q: How should each prisoner act?

- Betray is always the best action
- But,  $(-1, -1)$  is a better outcome for both
- Why? What goes wrong?
  - Selfish behaviors lead to inefficient outcome

# Example II: Markets on Amazon

Screenshot of an Amazon product page for "Artificial Intelligence: A Modern Approach (3rd Edition) (Hardcover)".

The page shows two listing options:

Price + Shipping	Condition (Learn more)	Delivery	Seller Information	Buying Options
<b>\$184.87</b> & FREE Shipping + \$0.00 estimated tax	New	<ul style="list-style-type: none"><li>Arrives between December 6 - 18.</li><li>Ships from CO, United States.</li><li><a href="#">Shipping rates and return policy.</a></li></ul>	<b>RushLtd</b> 95% positive over the past 12 months. (12,915 total ratings)	
<b>\$181.13</b>	New	<ul style="list-style-type: none"><li>Arrives between Nov. 29 -</li></ul>	<b>SuperBookDeal</b>	

Left sidebar filters:

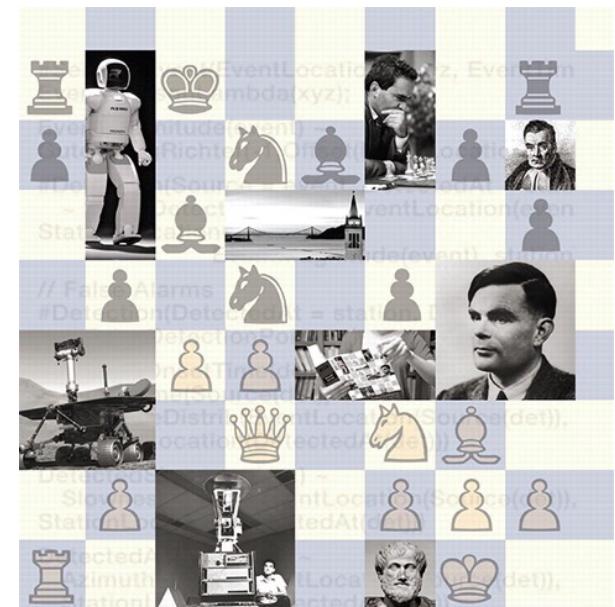
- Refine by [Clear all](#)
- Shipping**  
 ✓prime  
 Free shipping
- Condition**  
 New  
 Rental  
 Used

# Example II: Markets on Amazon

- Assume people will buy if the book price  $\leq \$200$
  - Product cost = \$20

If the market has only one book seller...

**Q: What price should this monopoly set?**



# Artificial Intelligence

## A Modern Approach

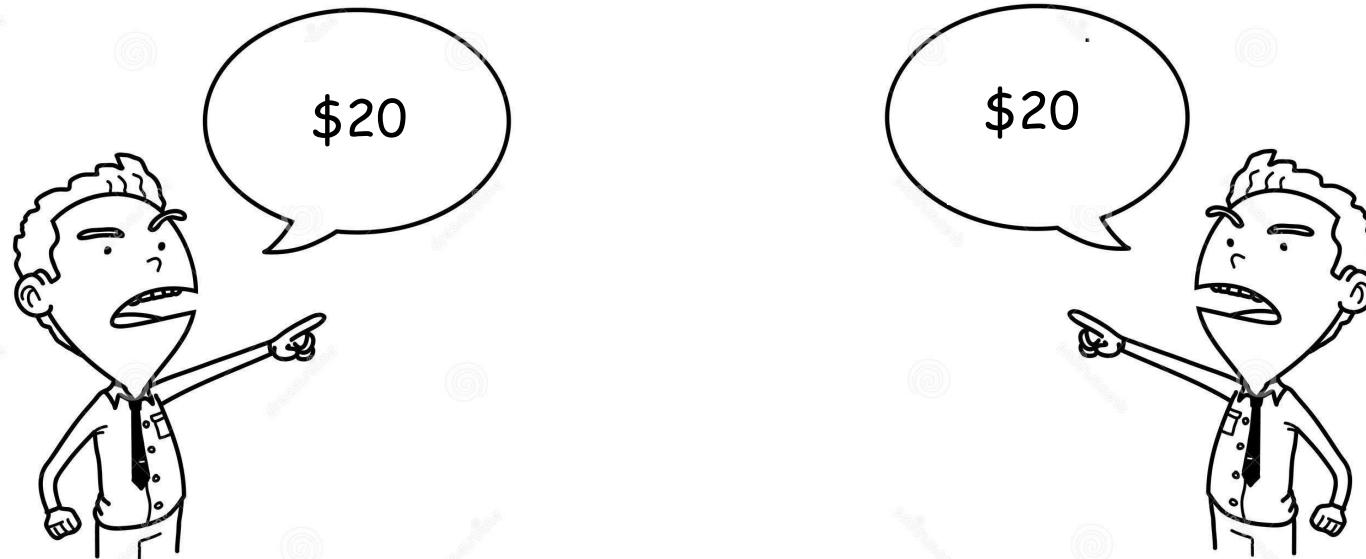
*Third Edition*

## Example II: Markets on Amazon

- Assume people will buy if the book price  $\leq \$200$
- Product cost = \$20

What if the market has **two** book sellers...

Q: What price should each seller set?



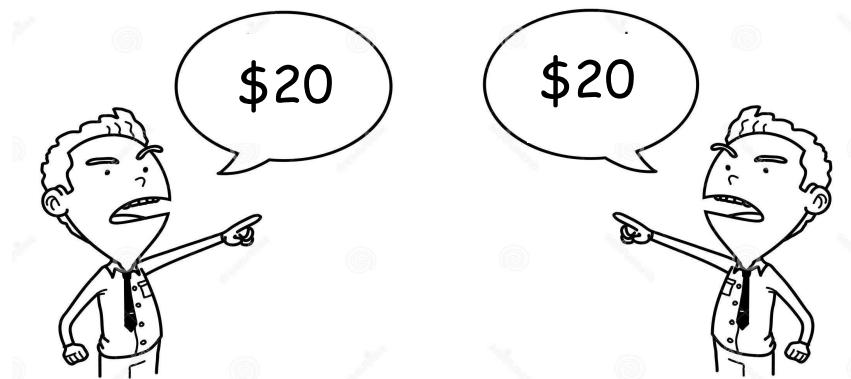
## Example II: Markets on Amazon

- Assume people will buy if the book price  $\leq \$200$
- Product cost = \$20

What if the market has **two** book sellers...

Q: What price should each seller set?

- The market reaches a “stable status” (a.k.a., equilibrium)
- Nobody can benefit via *unilateral deviation*



- Bertrand competition
- Seller's revenue-maximizing behaviors lead to low revenue

# Game Theory

Game Theory studies multiple-agent decision making in competitive scenarios where an agent's payoff depends on other agents' actions.

- Fundamental concept --- Equilibrium
  - A “stable status” at which any agent cannot improve his payoff through unilateral deviation
  - If exists, it is the solution concept (i.e., outcome) that we expect to happen
  - Resembles “optimal decision” in single-agent case
- A central theme in game theory is to study the equilibrium
  - Different “types” of equilibria
  - May not exist; even exist, not necessarily unique
  - Understand properties of equilibrium, compute equilibria, how to improve inefficiency of equilibrium . . .

# Machine Learning

- Difficult to give a universal definition
- At a high level, the task is to learn a function  $f: X \rightarrow Y$ , where  $(x, y) \in X \times Y$  is drawn from some distribution  $D$ 
  - **Input:** a set of samples  $\{(x_i, y_i)\}_{i=1,2,\dots,n}$  drawn from  $D$
  - **Output:** an algorithm  $A: X \rightarrow Y$  such that  $A(x) \approx f(x)$  (usually measured by some loss function)
- Examples
  - Classification:  $X$  = feature vectors;  $Y = \{0,1\}$
  - Regression:  $X$  = feature vectors;  $Y = \mathbb{R}$
  - Reinforcement learning has a slightly different setup, but can be thought as  $X$  = state space,  $Y$  = action space

# Problems at Interface of Learning and Game Theory

- If a game is unknown or too complex, can players learn to play the game optimally?
  - Yes, sometimes – no regret learning and convergence to equilibrium
- Can game-theoretic models inspire machine learning models?
  - Yes, GANs which are zero-sum games
- Data is the fuel for ML – Can we collect high-quality data from crowd?
  - Yes, via information elicitation mechanisms
- We know how to learn to recognize faces or languages, but can we also learn the design of games to achieve some goal?
  - Yes, learning optimal auctions, product pricing schemes, etc
- Gaming/strategic behaviors in ML? How to handle them?
  - Yes, e.g, learn whether to give loans to someone or whether to admit a student to UChicago based on their features
- ...

# Goodhart's Law

*When a measure becomes a target, it ceases to be a good measure*

# Main Topics of This Course

First Half: Machine learning for game theory

- Basics of linear programming and game theory
- No regret learning and its convergence to equilibrium

Second Half: Game theory for machine learning

- Incentivize high-quality data via information elicitation (a.k.a., crowdsourcing)
- Handle gaming behaviors in machine learning
  - Particularly, learning from strategic data sources, and fairness

# Main Topics of This Course

First Half: Machine learning for game theory

- Basics of linear programming and game theory
- No regret learning and its convergence to equilibrium

Second Half: Game theory for machine learning

- Incentivize high-quality data via information elicitation (a.k.a., crowdsourcing)
- Handle gaming behaviors in machine learning
  - Particularly, learning from strategic data sources, and fairness

Only cover fundamentals of each direction

# Main Topics of This Course

1 (Sep 27)	Introduction	Kleinberg/Leighton paper
2 (Sep 29)	Basics of LPs	Chapter 2.1, 2.2, 4.3 of <a href="#">Convex Optimization</a> by Boyd and Vandenberghe
3 (Oct 4)	LP duality	Lecture notes 5 and 6 of <a href="#">an optimization course by Trevisan</a>
4 (Oct 6)	Intro to Game Theory (I)	Section 3.1, 3.2, 3.3 of <a href="#">an game theory book by Shoham and Leyton-Brown</a>
5 (Oct 11)	Intro to Game Theory (II)	<a href="#">Equilibrium analysis of GANs</a> by Arora et al.
6 (Oct 13)	Intro to Online Learning	
7 (Oct 18)	Multiplicative Weight	<a href="#">A survey paper on MWU and its applications</a> by Arora et al.
8 (Oct 20)	Swap Regret	<a href="#">A note by Balcan on converting regret to swap regret</a>
9 (Oct 25)	Multi-Armed Bandits	Section 2 and 3 of the <a href="#">Book by Bubeck and Cesa-Bianchi on Bandits</a>
10 (Oct 27)	Prediction Markets (PMs)	<a href="#">Notes from a similar course by Waggoner</a>
11 (Nov 1)	PMs and Scoring Rules	<a href="#">The original paper</a> including all presented results
12 (Nov 3)	Peer Prediction	<a href="#">Bayesian Truth Serum paper</a>
13 (Nov 8)	Bayesian Persuasion	The original <a href="#">Bayesian Persuasion</a> paper
14 (Nov 10)	Pricing of Information	The <a href="#">Selling Information Through Consulting</a> paper
15 (Nov 15)	Strategic Learning I	<a href="#">PAC-learning for Strategic Classification</a> paper
16 (Nov 17)	Strategic Learning II	<a href="#">How Can ML Induce Right Efforts</a> paper
(Nov 22)	Thanksgiving, No Lecture	
(Nov 24)	Thanksgiving, No Lecture	
17 (Nov 29)	Tradeoffs of Fairness	<a href="#">Inherent Trade-Offs in the Fair Determination</a>
(Dec 1)	Project presentations	

# Course Goal

- Get familiar with basics of game theory and learning
- Understand machine learning questions in game-theoretic settings, and how to deal with some of them
- Understand gaming behaviors in machine learning applications, and how to deal with some of them
- Can understand cutting-edge research papers in relevant areas

# Targeted Audience of This Course

- Anyone planning to do research at the interface of game theory (or algorithm design) and machine learning
  - This is a new research direction with many opportunities/challenges
  - Recent breakthrough in no-limit poker is an example



# Targeted Audience of This Course

- Anyone planning to do research at the interface of game theory (or algorithm design) and machine learning
  - This is a new research direction with many opportunities/challenges
  - Recent breakthrough in no-limit poker is an example
- Anyone interested in theoretical ML, game theory, human factors in learning, AI
  - As more and more ML systems interact with human beings, such game-theoretic reasoning becomes increasingly important
  - With more techniques developed for ML, they also broadened our toolkits for designing and solving games
- Anyone interested in understanding basics of game theory and learning

# Who May not Be Suitable for This Course?

- Those who do not satisfy the prerequisites “in practice”
- Those who are looking for a recipe to implement ML/DL algorithms, or want to learn how to use TensorFlow, PyTorch, etc.
  - **This is primarily a theory course**
  - We will mostly focus on simple/basic yet theoretically insightful problems
  - The course is proof based – we will not write code

# Outline

- Course Overview
- Administrivia
- An Example

# Basic Information

- Course time: Tuesday/Thursday, 2:00 pm – 3:20 pm
- Lecture: in person (unless further university guidance is given)
- Instructor: Haifeng Xu
  - Email: [haifengxu@uchicago.edu](mailto:haifengxu@uchicago.edu)
  - Office Hour: **TBD**
- TAs
  - **Minbiao Han**: office hour **Wed/Fri 1 – 2 pm**, in person (room TBD)
- Depending on demand, can add more office hours (let us know!)
- Course website: [www.haifeng-xu.com/cmsc35401fa22/index.htm](http://www.haifeng-xu.com/cmsc35401fa22/index.htm)
  - Easier way is to search my personal website and navigate to course
- References: linked papers/notes on website, no official textbooks
  - Slides will be posted *after* lecture

# Prerequisites

- Mathematically mature: be comfortable with proofs
- Sufficient exposures to algorithms/optimization
  - CMSC 27200/27220 and equivalent
  - We will cover some basics of optimization

# Requirements and Grading

- 3 homeworks, 50% of grade.
  - Proof based, and will be challenging
  - Discussion allowed, even encouraged, but must write up solutions independently
  - Must be written up in Latex – hand-written solutions will not be accepted
  - One late homework allowed, at most 2 days
- Research project, 50% of grade. Project instructions will be posted on website later.
  - Team up: 2 – 4 people per team
  - Can thoroughly survey a research field, or
  - Study a relevant research question, e.g., arising from your own research
  - Presentation form: a report in PDF
- FYI: no need to worry about your grade if you do invest time

If you have any suggestions/comments/concerns,  
feel free to email me.

# Outline

- Course Overview
- Administrivia
- An Example

# Learning to Sell a Product

- You are a product seller facing  $N$  unknown buyers
- These buyers all value your product at the same  $\nu \in [0,1]$ , which however is *unknown* to you
- Buyers come in sequence  $1, 2, \dots, N$ ; For each buyer, you can choose a price  $p$  and ask him whether he is willing to buy the product
  - If  $\nu \geq p$ , she/he purchases; otherwise leaves the queue



# Learning to Sell a Product

- You are a product seller facing  $N$  unknown buyers
- These buyers all value your product at the same  $\nu \in [0,1]$ , which however is *unknown* to you
- Buyers come in sequence  $1, 2, \dots, N$ ; For each buyer, you can choose a price  $p$  and ask him whether he is willing to buy the product
  - If  $\nu \geq p$ , she/he purchases; otherwise not
- How to quickly learn these buyers' value  $\nu$  within precision  $\epsilon = \frac{1}{N}$ ?
  - This is a pure learning problem
  - (Well, you may directly ask a buyer's value, but guess what will happen?)
- Answer:  $\log(N)$  rounds via BinarySearch

# Learning to Sell a Product

- You are a product seller facing  $N$  unknown buyers
- These buyers all value your product at the same  $\nu \in [0,1]$ , which however is *unknown* to you

Let us move to a natural game-theoretic setup .....

- You have an ultimate objective of maximizing your revenue, but do not really care about learning  $\nu$  (though you may have to)
- How much revenue can BinarySearch secure?
  - May get really unlucky in first  $\log(N)$  rounds and no sale happened
  - After  $\log(N)$  rounds, can set a price  $p \geq \tilde{\nu} - 1/N$  ( $\tilde{\nu}$  is learned value)

$$\text{Rev} = \underbrace{0}_{\text{First } \log(N) \text{ rounds}} + \underbrace{(N - \log N)(\nu - \frac{2}{N})}_{\text{Remaining rounds}} \approx \nu N - \nu \log N - 2$$

# Regret as Performance Measure

- To measure algorithm performance, we use **regret**

**Regret** := how much less is an algorithm's utility compared to the (idealized) case where we know  $\nu$ .

- Had we known  $\nu$ , should just price the product at  $p = \nu$ , earning  $\nu N$
- The regret is then

$$\text{Regret}(\text{binary search}) \approx \nu N - [\nu N - \nu \log N - 2] = \nu \log N + 2$$

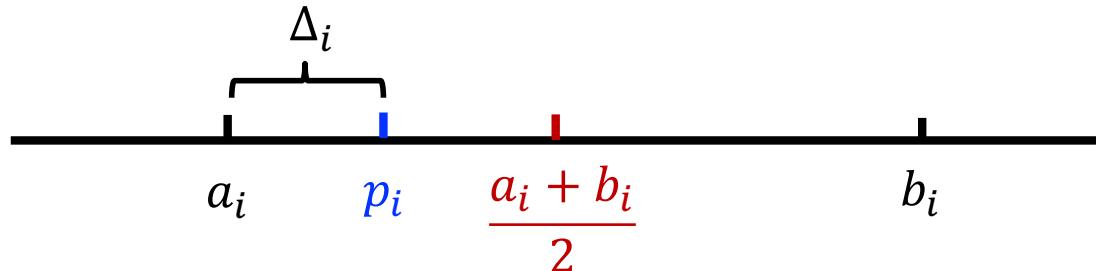
Q: Is this the best (i.e., the smallest) regret?

# An Algorithm with Smaller Regret

**Theorem [Kleinberg/Leighton, FOCS'03]** : there is an algorithm achieving regret at most  $(1 + 2 \log \log N)$

Why BinarySearch may be bad?

- For buyer  $i$ , BinarySearch maintains an interval bound  $[a_i, b_i]$  and use  $p_i = (a_i + b_i)/2$  for buyer  $i$ 
  - This learns  $v$  as quickly as possible
  - But maybe bad for revenue since we will get 0 revenue if  $p_i > v$ , and  $p_i = (a_i + b_i)/2$  may be too high/aggressive

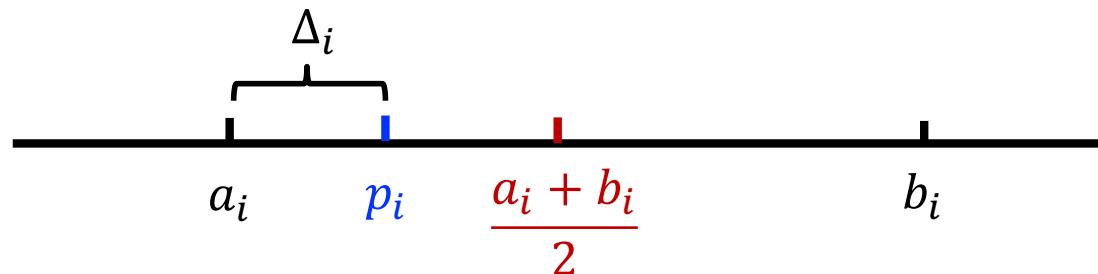


# An Algorithm with Smaller Regret

**Theorem [Kleinberg/Leighton, FOCS'03]** : there is an algorithm achieving regret at most  $(1 + 2 \log \log N)$

Why BinarySearch may be bad?

- For buyer  $i$ , BinarySearch maintains an interval bound  $[a_i, b_i]$  and use  $p_i = (a_i + b_i)/2$  for buyer  $i$ 
  - This learns  $v$  as quickly as possible
  - But maybe bad for revenue since we will get 0 revenue if  $p_i > v$ , and  $p_i = (a_i + b_i)/2$  may be too high/aggressive
- Algorithm idea: use more conservative prices

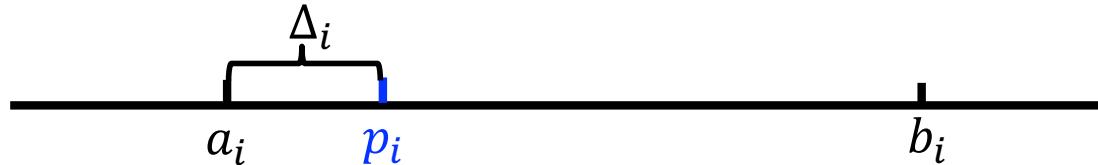


# An Algorithm with Smaller Regret

**Theorem [Kleinberg/Leighton, FOCS'03]** : there is an algorithm achieving regret at most  $(1 + 2 \log \log N)$

The Algorithm (note  $v \in [0,1]$ ):

- Maintains an interval bound  $[a_i, b_i]$  and a **step size  $\Delta_i$**
- Offer price  $p_i = a_i + \Delta_i$  for buyer  $i$



- If  $i$  accepts, update  $a_{i+1} = p_i, b_{i+1} = b_i, \Delta_{i+1} = \Delta_i$

# An Algorithm with Smaller Regret

**Theorem [Kleinberg/Leighton, FOCS'03]** : there is an algorithm achieving regret at most  $(1 + 2 \log \log N)$

The Algorithm (note  $v \in [0,1]$ ):

- Maintains an interval bound  $[a_i, b_i]$  and a **step size**  $\Delta_i$
- Offer price  $p_i = a_i + \Delta_i$  for buyer  $i$



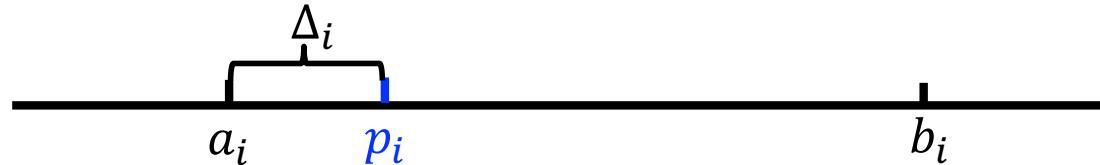
- If  $i$  accepts, update  $a_{i+1} = p_i$ ,  $b_{i+1} = b_i$ ,  $\Delta_{i+1} = \Delta_i$

# An Algorithm with Smaller Regret

**Theorem [Kleinberg/Leighton, FOCS'03]** : there is an algorithm achieving regret at most  $(1 + 2 \log \log N)$

The Algorithm (note  $v \in [0,1]$ ):

- Maintains an interval bound  $[a_i, b_i]$  and a step size  $\Delta_i$
- Offer price  $p_i = a_i + \Delta_i$  for buyer  $i$



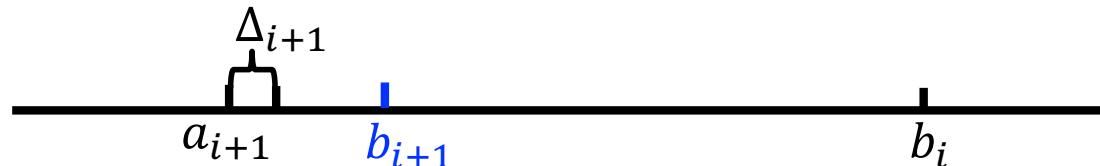
- If  $i$  accepts, update  $a_{i+1} = p_i, b_{i+1} = b_i, \Delta_{i+1} = \Delta_i$
- Otherwise, update  $a_{i+1} = a_i, b_{i+1} = p_i, \Delta_{i+1} = (\Delta_i)^2$

# An Algorithm with Smaller Regret

**Theorem [Kleinberg/Leighton, FOCS'03]** : there is an algorithm achieving regret at most  $(1 + 2 \log \log N)$

The Algorithm (note  $v \in [0,1]$ ):

- Maintains an interval bound  $[a_i, b_i]$  and a step size  $\Delta_i$
- Offer price  $p_i = a_i + \Delta_i$  for buyer  $i$



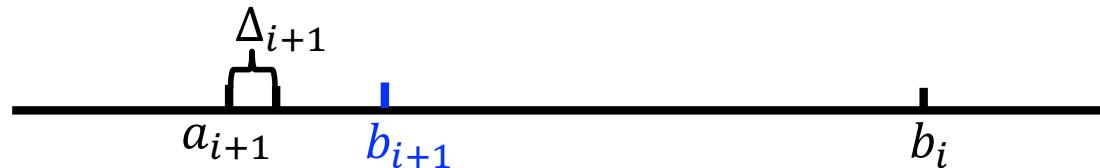
- If  $i$  accepts, update  $a_{i+1} = p_i$ ,  $b_{i+1} = b_i$ ,  $\Delta_{i+1} = \Delta_i$
- Otherwise, update  $a_{i+1} = a_i$ ,  $b_{i+1} = p_i$ ,  $\Delta_{i+1} = (\Delta_i)^2$

# An Algorithm with Smaller Regret

**Theorem [Kleinberg/Leighton, FOCS'03]** : there is an algorithm achieving regret at most  $(1 + 2 \log \log N)$

The Algorithm (note  $v \in [0,1]$ ):

- Maintains an interval bound  $[a_i, b_i]$  and a step size  $\Delta_i$
- Offer price  $p_i = a_i + \Delta_i$  for buyer  $i$



- If  $i$  accepts, update  $a_{i+1} = p_i$ ,  $b_{i+1} = b_i$ ,  $\Delta_{i+1} = \Delta_i$
- Otherwise, update  $a_{i+1} = a_i$ ,  $b_{i+1} = p_i$ ,  $\Delta_{i+1} = (\Delta_i)^2$
- Start with  $a_1 = 0$ ,  $b_1 = 1$ ,  $\Delta_1 = 1/2$ ; Once  $b_i - a_i \leq \frac{1}{N}$ , always use  $p = a_i$  afterwards

Remark: searching smaller region with smaller step size.

# An Algorithm with Smaller Regret

**Theorem [Kleinberg/Leighton, FOCS'03]** : there is an algorithm achieving regret at most  $(1 + 2 \log \log N)$

Algorithm analysis:



**Claim 1:** The step size  $\Delta_i$  takes values  $2^{-2^j}$  for  $j = 0, 1, \dots$ .

Moreover, whenever  $\Delta_{i+1} = (\Delta_i)^2$  happens,  $b_{i+1} - a_{i+1} = \sqrt{\Delta_{i+1}}$ .

Proof

- Recall  $\Delta_1 = \frac{1}{2} = 2^{-2^0}$ , and step size update  $\Delta_{i+1} = (\Delta_i)^2$
- If  $\Delta_i = 2^{-2^j}$ , then  $(\Delta_i)^2 = 2^{-2^{j+1}} = 2^{-2^{j+1}}$
- When  $\Delta_{i+1} = (\Delta_i)^2$  happens,  $b_{i+1} - a_{i+1} = \Delta_i = \sqrt{\Delta_{i+1}}$

# An Algorithm with Smaller Regret

**Theorem [Kleinberg/Leighton, FOCS'03]** : there is an algorithm achieving regret at most  $(1 + 2 \log \log N)$

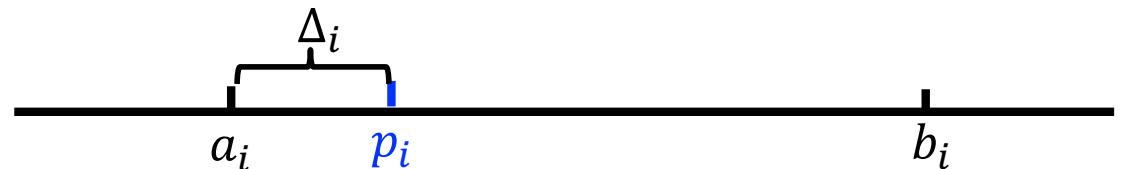
Algorithm analysis:



- After  $b_i - a_i \leq \frac{1}{N}$ , the total regret is at most 1
  - Because (1) regret of each step is at most  $\frac{1}{N}$ ; (2) there are at most  $N$  rounds
- Main step is to bound regret before reaching  $b_i - a_i = \frac{1}{N}$

# An Algorithm with Smaller Regret

**Theorem [Kleinberg/Leighton, FOCS'03]** : there is an algorithm achieving regret at most  $(1 + 2 \log \log N)$



Algorithm analysis:

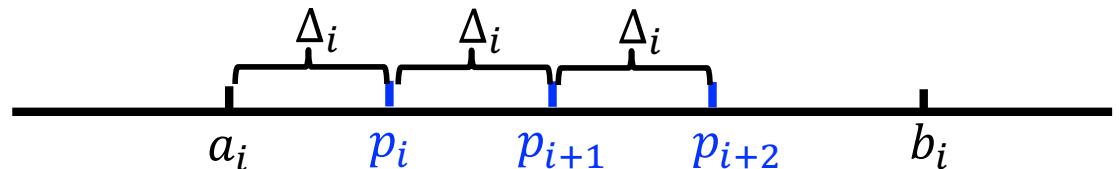
- How many **step size value updates** needed to reach  $b_i - a_i = \frac{1}{N}$ ?
  - **log log N**: set  $2^{-2^i} = \frac{1}{N} \rightarrow i = \log \log N$
  - The following claim then completes the proof of the theorem

**Claim 2:** total regret from any **step size value**  $\Delta$  is at most 2.

- No sale happens only once for any step size  $\rightarrow$  regret at most 1

# An Algorithm with Smaller Regret

**Theorem [Kleinberg/Leighton, FOCS'03]** : there is an algorithm achieving regret at most  $(1 + 2 \log \log N)$



Algorithm analysis:

- How many **step size value updates** needed to reach  $b_i - a_i = \frac{1}{N}$ ?
  - **log log N**: set  $2^{-2^i} = \frac{1}{N} \rightarrow i = \log \log N$
  - The following claim then completes the proof of the theorem

**Claim 2:** total regret from any **step size value**  $\Delta$  is at most 2.

- No sale happens only once for any step size  $\rightarrow$  regret at most 1
- What about the regret when sales happen?
  - Can happen at most  $\sqrt{\Delta}/\Delta$  times since  $b_i - a_i \leq \sqrt{\Delta}$ ; regret from each time is at most  $b_i - a_i (\leq \sqrt{\Delta})$
  - Regret from sales is at most  $(\sqrt{\Delta}/\Delta) \times \sqrt{\Delta} = 1$

# An Algorithm with Smaller Regret

## Remarks

- $O(\log \log N)$  is also the order-wise best regret [KL, FOCS'13]
- This is an example of **exploration** vs **exploitation**
  - Exploration: want to learn  $\nu$
  - Exploitation: but ultimate goal is to utilize learned  $\nu$  to maximize revenue
  - More in later lectures...
- BinarySearch is best for exploration, but did not balance the two
- The “optimal” algorithm uses less step value updates, but more interval updates
  - Less step value updates are to be conservative about prices in order for revenue maximization
  - More interval updates mean interacting with more buyers to learn  $\nu$
  - That is, **slower learning** but **higher revenue**

# Well, This is Not the End Yet . . .

- Here, it is crucial that each buyer only shows up once
- What if the same buyer shows up repeatedly?
  - In fact, this is more realistic
  - E.g., in online advertising, buyer = an advertiser
- How should a (repeatedly showing up) buyer behave if he knows seller is learning her value  $v$  and then uses it to set a price for her?

## Open Research Questions:

1. How to design pricing schemes for a repeatedly showing up buyer to maximize revenue when the buyer knows you are learning his value?
2. How to generalize to selling multiple products?

# Thank You

Haifeng Xu

University of Chicago

[haifengxu@uchicago.edu](mailto:haifengxu@uchicago.edu)