

# Announcements

- HW 1 draft is slightly updated; See website for more info
- Minbiao's office hour has been changed to **Thursday 1-2 pm** from this week, at **Rice Hall 442**

# CS6501:Topics in Learning and Game Theory (Fall 2019)

## MW Updates and Implications

---

---

Instructor: Haifeng Xu

# Outline

- Regret Proof of MW Update
- Convergence to Minimax Equilibrium
- Convergence to Coarse Correlated Equilibrium

# Recap: the Model of Online Learning

At each time step  $t = 1, \dots, T$ , the following occurs in order:

1. Learner picks a distribution  $p_t$  over actions  $[n]$
2. Adversary picks cost vector  $c_t \in [0,1]^n$
3. Action  $i_t \sim p_t$  is chosen and learner incurs cost  $c_t(i_t)$
4. Learner observes  $c_t$  (for use in future time steps)

- Learner's goal: pick distribution sequence  $p_1, \dots, p_T$  to minimize expected cost  $\mathbb{E} [\sum_{t \in T} c_t(i_t)]$
- Expectation over randomness of action

# Measure Algorithms via Regret

- Regret – how much the learner regrets, had he known the cost vector  $c_1, \dots, c_T$  in hindsight
- Formally,

$$R_T = \mathbb{E}_{i_t \sim p_t} \sum_{t \in [T]} c_t(i_t) - \min_{i \in [n]} \sum_{t \in [T]} c_t(i)$$

- Benchmark  $\min_{i \in [n]} \sum_t c_t(i)$  is the learner utility had he known  $c_1, \dots, c_T$  and is allowed to take the best **single action across all rounds**
  - Can also use other benchmarks, but  $\min_{i \in [n]} \sum_t c_t(i)$  is mostly used

An algorithm has **no regret** if  $\frac{R_T}{T} \rightarrow 0$  as  $T \rightarrow \infty$ , i.e.,  $R_T = o(T)$ .

Regret is an appropriate performance measure of online algorithms

- It measures exactly the loss due to not knowing the data in advance

# The Multiplicative Weight Update Alg

Parameter:  $\epsilon$

Initialize weight  $w_1(i) = 1, \forall i = 1, \dots, n$

For  $t = 1, \dots, T$

1. Let  $W_t = \sum_{i \in [n]} w_t(i)$ , pick action  $i$  with probability  $w_t(i)/W_t$
2. Observe cost vector  $c_t \in [0,1]^n$
3. For all  $i \in [n]$ , update  $w_{t+1}(i) = w_t(i) \cdot (1 - \epsilon \cdot c_t(i))$

**Theorem.** MW Update algorithm achieves regret at most  $O(\sqrt{T \ln n})$  for the previously described online learning problem.

- Last lecture: both  $\sqrt{T}$  and  $\ln n$  term are necessary
- Next, we prove the theorem

# Intuition of the Proof

Parameter:  $\epsilon$

Initialize weight  $w_1(i) = 1, \forall i = 1, \dots, n$

For  $t = 1, \dots, T$

1. Let  $W_t = \sum_{i \in [n]} w_t(i)$ , pick action  $i$  with probability  $w_t(i)/W_t$
2. Observe cost vector  $c_t \in [0,1]^n$
3. For all  $i \in [n]$ , update  $w_{t+1}(i) = w_t(i) \cdot (1 - \epsilon \cdot c_t(i))$

➤ The decrease of weights relates to expected cost at each round

- Expected cost at round  $t$  is  $\bar{C}_t = \sum_{i \in [n]} p_t(i) \cdot c_t(i) = \frac{\sum_{i \in [n]} w_t(i) \cdot c_t(i)}{W_t}$
- Propositional to the **decrease of total weight** at round  $t$ , which is

$$\sum_{i \in [n]} \epsilon \cdot w_t(i) c_t(i) = \epsilon W_t \cdot \bar{C}_t$$

➤ Proof idea: bound how fast do total weights decrease

# Proof Step I: How Fast do Total Weights Decrease?

**Lemma 1.**  $W_{t+1} \leq W_t \cdot e^{-\epsilon \bar{C}_t}$  where  $W_t = \sum_{i \in [n]} w_t(i)$  is the total weight at  $t$  and  $\bar{C}_t$  is the expected loss at time  $t$ .

$$\bar{C}_t = \sum_{i \in [n]} p_t(i) c_t(i) = \frac{\sum_{i \in [n]} w_t(i) c_t(i)}{W_t}$$

Proof

➤ Almost Immediate from update rule  $w_{t+1}(i) = w_t(i) \cdot (1 - \epsilon \cdot c_t(i))$

$$\begin{aligned} W_{t+1} &= \sum_{i \in [n]} w_{t+1}(i) \\ &= \sum_{i \in [n]} w_t(i) \cdot (1 - \epsilon \cdot c_t(i)) \\ &= W_t - \epsilon \cdot \sum_{i \in [n]} w_t(i) \cdot c_t(i) \\ &= W_t - \epsilon \cdot W_t \bar{C}_t = W_t(1 - \epsilon \cdot \bar{C}_t) \\ &\leq W_t \cdot e^{-\epsilon \cdot \bar{C}_t} \quad \text{since } 1 - \delta \leq e^{-\delta}, \forall \delta \geq 0 \end{aligned}$$

# Proof Step I: How Fast do Total Weights Decrease?

**Lemma 1.**  $W_{t+1} \leq W_t \cdot e^{-\epsilon \bar{C}_t}$  where  $W_t = \sum_{i \in [n]} w_t(i)$  is the total weight at  $t$  and  $\bar{C}_t$  is the expected loss at time  $t$ .

$$\bar{C}_t = \sum_{i \in [n]} p_t(i) c_t(i) = \frac{\sum_{i \in [n]} w_t(i) c_t(i)}{W_t}$$

**Corollary 1.**  $W_{T+1} \leq n e^{-\epsilon \sum_{t=1}^T \bar{C}_t}$ .

$$\begin{aligned} W_{T+1} &\leq W_T \cdot e^{-\epsilon \bar{C}_T} \\ &\leq [W_{T-1} \cdot e^{-\epsilon \bar{C}_{T-1}}] \cdot e^{-\epsilon \bar{C}_T} \\ &= W_{T-1} \cdot e^{-\epsilon [\bar{C}_T + \bar{C}_{T-1}]} \\ &\quad \dots \\ &= W_1 \cdot e^{-\epsilon \cdot \sum_{t=1}^T \bar{C}_t} \\ &= n \cdot e^{-\epsilon \cdot \sum_{t=1}^T \bar{C}_t} \end{aligned}$$

## Proof Step 2: Lower Bounding $W_{T+1}$

**Lemma 2.**  $W_{T+1} \geq e^{-T\epsilon^2} \cdot e^{-\epsilon \sum_{t=1}^T c_t(i)}$  for any action  $i$ .

$$\begin{aligned} W_{T+1} &\geq w_{T+1}(i) \\ &= w_1(i)(1 - \epsilon c_1(i))(1 - \epsilon c_2(i)) \dots (1 - \epsilon c_T(i)) \quad \text{by MW update rule} \\ &\geq \prod_{t=1}^T e^{-\epsilon c_t(i) - \epsilon^2 [c_t(i)]^2} \quad \text{by fact } 1 - \delta \geq e^{-\delta - \delta^2} \\ &\geq e^{-T\epsilon^2} \cdot e^{-\epsilon \sum_{t=1}^T c_t(i)} \quad \text{relax } [c_t(i)]^2 \text{ to 1} \end{aligned}$$

# Proof Step 3: Combing the Two Lemmas

**Corollary 1.**  $W_{T+1} \leq n e^{-\epsilon \sum_{t=1}^T \bar{c}_t}$ .

**Lemma 2.**  $W_{T+1} \geq e^{-T\epsilon^2} \cdot e^{-\epsilon \sum_{t=1}^T c_t(i)}$  for any action  $i$ .

➤ Therefore, for any  $i$  we have

$$\begin{aligned} & e^{-T\epsilon^2} \cdot e^{-\epsilon \sum_{t=1}^T c_t(i)} \leq n e^{-\epsilon \sum_{t=1}^T \bar{c}_t} \\ \Leftrightarrow & -T\epsilon^2 - \epsilon \sum_{t=1}^T c_t(i) \leq \ln n - \epsilon \sum_{t=1}^T \bar{c}_t && \text{take “ln” on both sides} \\ \Leftrightarrow & \sum_{t=1}^T \bar{c}_t - \sum_{t=1}^T c_t(i) \leq \frac{\ln n}{\epsilon} + T\epsilon && \text{rearrange terms} \end{aligned}$$

Taking  $\epsilon = \sqrt{\ln n / T}$ , we have

$$\sum_{t=1}^T \bar{c}_t - \min_i \sum_{t=1}^T c_t(i) \leq 2\sqrt{T \ln n}$$

# Remarks

- Some MW description uses  $w_{t+1}(i) = w_t(i) \cdot e^{-\epsilon \cdot c_t(i)}$ . Analysis is similar due to the fact  $e^{-\epsilon} \approx 1 - \epsilon$  for small  $\epsilon \in [0,1]$
- The same algorithm also works for  $c_t \in [-\rho, \rho]$  (still use update rule  $w_{t+1}(i) = w_t(i) \cdot (1 - \epsilon \cdot c_t(i))$ ). Analysis is the same
- MW update is a very powerful technique – it can also be used to solve, e.g., LP, semidefinite programs, SetCover, Boosting, etc.
  - Because it works for **arbitrary cost vectors**
  - Next, we show how it can be used to compute equilibria of games where the “cost vector” will be generated by other players

# Outline

- Regret Proof of MW Update
- Convergence to Minimax Equilibrium
- Convergence to Coarse Correlated Equilibrium

## Online learning – A natural way to play [repeated games](#)

Repeated game: the same game played for many rounds

- Think about how you play rock-paper-scissor repeatedly
- In reality, we play like online learning
  - You try to analyze the past patterns, then decide which action to respond, possibly with some randomness
  - This is basically online learning!



# Repeated Zero-Sum Games with No-Regret Players

Basic Setup:

- A zero-sum game with payoff matrix  $U \in \mathbb{R}^{m \times n}$
- Row player maximizes utility and has actions  $[m] = \{1, \dots, m\}$ 
  - Column player thus minimizes utility
- The game is played **repeatedly** for  $T$  rounds
- Each player uses an online learning algorithm to pick a mixed strategy at each round

# Repeated Zero-Sum Games with No-Regret Players

- From row player's perspective, the following occurs in order at round  $t$ 
  - Picks a mixed strategy  $x_t \in \Delta_m$  over actions in  $[m]$
  - Her opponent, the column player, picks a mixed strategy  $y_t \in \Delta_n$
  - Action  $i_t \sim x_t$  is chosen and row player receives utility  $U(i_t, y_t) = \sum_{j \in [n]} y_t(j) \cdot U(i_t, j)$
  - Row player learns  $y_t$  (for future use)
- Column player has a symmetric perspective, but will think of  $U(i, j)$  as his cost

Difference from online learning: utility/cost vector determined by the opponent, instead of being arbitrarily chosen

# Repeated Zero-Sum Games with No-Regret Players

- Expected total utility of row player  $\sum_{t=1}^T U(x_t, y_t)$ 
  - Note:  $U(x_t, y_t) = \sum_{i,j} U(i, j)x_t(i)y_t(j) = (x_t)^T U y_t$

- Regret of row player is

$$\max_{i \in [m]} \sum_{t=1}^T U(i, y_t) - \sum_{t=1}^T U(x_t, y_t)$$

- Regret of column player is

$$\sum_{t=1}^T U(x_t, y_t) - \min_{j \in [n]} \sum_{t=1}^T U(x_t, j)$$

# From No Regret to Minimax Theorem

Next, we give another proof of the minimax theorem, using the fact that no regret algorithms exist (e.g., MW update)

# From No Regret to Minimax Theorem

- Assume both players use no-regret learning algorithms
- For row player, we have

$$\begin{aligned} R_T^{row} &= \max_{i \in [m]} \sum_{t=1}^T U(i, y_t) - \sum_{t=1}^T U(x_t, y_t) \\ \Leftrightarrow \frac{1}{T} \sum_{t=1}^T U(x_t, y_t) + \frac{R_T^{row}}{T} &= \frac{1}{T} \max_{i \in [m]} \sum_{t=1}^T U(i, y_t) \\ &= \max_{i \in [m]} U\left(i, \frac{\sum_t y_t}{T}\right) \\ &\geq \min_{y \in \Delta_n} \max_{i \in [m]} U(i, y) \end{aligned}$$

# From No Regret to Minimax Theorem

- Assume both players use no-regret learning algorithms
- For row player, we have

$$\frac{1}{T} \sum_{t=1}^T U(x_t, y_t) + \frac{R_T^{row}}{T} \geq \min_{y \in \Delta_n} \max_{i \in [m]} U(i, y)$$

- Similarly, for column player,

$$R_T^{column} = \sum_{t=1}^T U(x_t, y_t) - \min_{j \in [n]} \sum_{t=1}^T U(x_t, j)$$

implies

$$\frac{1}{T} \sum_{t=1}^T U(x_t, y_t) - \frac{R_T^{column}}{T} \leq \max_{x \in \Delta_m} \min_{j \in [n]} U(x, j)$$

- Let  $T \rightarrow \infty$ , no regret implies  $\frac{R_T^{row}}{T}$  and  $\frac{R_T^{column}}{T}$  tend to 0. We have

$$\min_{y \in \Delta_n} \max_{i \in [m]} U(i, y) \leq \max_{x \in \Delta_m} \min_{j \in [n]} U(x, j)$$

# From No Regret to Minimax Theorem

- Assume both players use no-regret learning algorithms

$$\frac{1}{T} \sum_{t=1}^T U(x_t, y_t) + \frac{R_T^{row}}{T} \geq \min_{y \in \Delta_n} \max_{i \in [m]} U(i, y)$$

$$\frac{1}{T} \sum_{t=1}^T U(x_t, y_t) - \frac{R_T^{column}}{T} \leq \max_{x \in \Delta_m} \min_{j \in [n]} U(x, j)$$

$$\Rightarrow \min_{y \in \Delta_n} \max_{i \in [m]} U(i, y) \leq \max_{x \in \Delta_m} \min_{j \in [n]} U(x, j)$$

- Recall that  $\text{min-max} \geq \text{max-min}$  also holds, because moving second will not be worse for the row player

**Corollary.**  $\frac{1}{T} \sum_{t=1}^T U(x_t, y_t)$  converges to the game value

# Convergence to Nash Equilibrium

**Theorem.** Suppose both players use no-regret learning algorithms with action sequence  $\{x_t\}$  and  $\{y_t\}$ . Then  $\frac{1}{T} \sum_{t=1}^T U(x_t, y_t)$  converges to the game value and  $(\frac{\sum_{t=1}^T x_t}{T}, \frac{\sum_{t=1}^T y_t}{T})$  converges to NE of the game.

- Recall that  $(x^*, y^*)$  is a NE if and only if  $x^*$  is the maximin strategy and  $y^*$  is the minimax strategy
- From previous derivations

$$\begin{aligned}\frac{1}{T} \sum_{t=1}^T U(x_t, y_t) + \frac{R_T^{row}}{T} &= \max_{i \in [m]} U\left(i, \frac{\sum_t y_t}{T}\right) \\ &\geq \min_{y \in \Delta_n} \max_{i \in [m]} U(i, y)\end{aligned}$$

- As  $T \rightarrow \infty$ , “ $\geq$ ” becomes “ $=$ ”. So  $\frac{\sum_t y_t}{T}$  solves the min-max problem
- Similarly,  $\frac{\sum_t x_t}{T}$  solves the max-min problem

# Remarks

- If both players use no regret algorithms with  $O(\sqrt{T})$ , then  $\frac{1}{T} \sum_{t=1}^T U(x_t, y_t)$  converges to the game value at rate  $\frac{R_T}{T} = \frac{1}{\sqrt{T}}$
- This convergence rate can be improved to  $\frac{1}{T}$  by careful regularization of the no-regret algorithm
  - More readings: “*Fast Convergence of Regularized Learning in Games*” [**NIPS’15 best paper**]
  - Intuition: our no-regret algorithm assumes adversarial feedbacks but the other player is not really adversary – he uses another no-regret algorithm
  - This can be exploited to improve learning rate

# Remarks

- Convergence of no-regret learning to NE is the key framework for designing the AI agent that beats top humans in Texas hold'em poker
  - Plus many other game solving techniques and engineering work
  - More reading: “Safe and Nested Subgame Solving for Imperfect-Information Games.” [NeurIPS’17 best paper]



# Remarks

- Convergence of no-regret learning to NE is the key framework for designing the AI agent that beats top humans in Texas hold'em poker
  - Plus many other game solving techniques and engineering work
  - More reading: “Safe and Nested Subgame Solving for Imperfect-Information Games.” [NeurIPS’17 best paper]

Exciting research is happening at this intersected space of Learning & Game Theory



# Outline

- Regret Proof of MW Update
- Convergence to Minimax Equilibrium
- Convergence to Coarse Correlated Equilibrium

# Recap: Normal-Form Games and CCE

- $n$  players, denoted by set  $[n] = \{1, \dots, n\}$
- Player  $i$  takes action  $a_i \in A_i$
- Player utility depends on the outcome of the game, i.e., an action profile  $a = (a_1, \dots, a_n)$ 
  - Player  $i$  receives payoff  $u_i(a)$  for any outcome  $a \in \prod_{i=1}^n A_i$
- Course correlated equilibrium is an action recommendation policy

A recommendation policy  $\pi$  is a **coarse correlated equilibrium** if

$$\sum_{a \in A} u_i(a) \cdot \pi(a) \geq \sum_{a \in A} u_i(\textcolor{red}{a'}_i, a_{-i}) \cdot \pi(a), \forall \textcolor{red}{a'}_i \in A_i, \forall i \in [n].$$

That is, for any player  $i$ , following  $\pi$ 's recommendations is better than opting out of the recommendation and “acting on his own”.

# Repeated Games with No-Regret Players

- The game is played repeatedly for  $T$  rounds
- Each player uses an online learning algorithm to select a mixed strategy at each round  $t$
- For any **player  $i$** 's perspective, the following occurs in order at  $t$ 
  - Picks a mixed strategy  $x_i^t \in \Delta_{|A_i|}$  over actions in  $A_i$
  - Any other player  $j \neq i$  picks a mixed strategy  $x_j^t \in \Delta_{|A_j|}$
  - Player  $i$  receives expected utility  $U_i(x_i^t, x_{-i}^t) = \mathbb{E}_{a \sim (x_i^t, x_{-i}^t)} u_i(a)$
  - Player  $i$  learns  $x_{-i}^t$  (for future use)

# Repeated Games with No-Regret Players

- Expected total utility of player  $i$  equals  $\sum_{t=1}^T U_i(x_i^t, x_{-i}^t)$
- Regret of player  $i$  is

$$R_i^T = \max_{a_i \in A_i} \sum_{t=1}^T U(a_i, x_{-i}^t) - \sum_{t=1}^T U_i(x_i^t, x_{-i}^t)$$

# From No Regret to CCE

**Theorem.** Suppose all players use no-regret learning algorithms with strategy sequence  $\{x_i^t\}_{t \in [T]}$  for  $i$ . The following recommendation policy  $\pi^T$  converges to a CCE:  $\pi^T(a) = \frac{1}{T} \sum_t \prod_{i \in [n]} x_i^t(a_i), \forall a \in A$ .

Remarks:

- In mixed strategy profile  $(x_1^t, x_2^t, \dots, x_n^t)$ , prob of  $a$  is  $\prod_{i \in [n]} x_i^t(a_i)$
- $\pi^T(a)$  is simply the average of  $\prod_{i \in [n]} x_i^t(a_i)$  over  $T$  rounds

# From No Regret to CCE

**Theorem.** Suppose all players use no-regret learning algorithms with strategy sequence  $\{x_i^t\}_{t \in [T]}$  for  $i$ . The following recommendation policy  $\pi^T$  converges to a CCE:  $\pi^T(a) = \frac{1}{T} \sum_t \Pi_{i \in [n]} x_i^t(a_i), \forall a \in A$ .

Remarks:

- In mixed strategy profile  $(x_1^t, x_2^t, \dots, x_n^t)$ , prob of  $a$  is  $\Pi_{i \in [n]} x_i^t(a_i)$
- $\pi^T(a)$  is simply the average of  $\Pi_{i \in [n]} x_i^t(a_i)$  over  $T$  rounds
- Player  $i$ 's expected utility from  $\pi^T$  is

$$\begin{aligned} & \sum_{a \in A} \left[ \frac{1}{T} \sum_t \Pi_{i \in [n]} x_i^t(a_i) \right] \cdot u_i(a) \\ &= \frac{1}{T} \sum_t \sum_{a \in A} \Pi_{i \in [n]} x_i^t(a_i) \cdot u_i(a) \\ &= \frac{1}{T} \sum_t u_i(x_i^t, x_{-i}^t) \end{aligned}$$

# From No Regret to CCE

**Theorem.** Suppose all players use no-regret learning algorithms with strategy sequence  $\{x_i^t\}_{t \in [T]}$  for  $i$ . The following recommendation policy  $\pi^T$  converges to a CCE:  $\pi^T(a) = \frac{1}{T} \sum_t \prod_{i \in [n]} x_i^t(a_i), \forall a \in A$ .

Proof:

- The CCE condition requires for all player  $i$

$$\frac{1}{T} \sum_t u_i(x_i^t, x_{-i}^t) \geq \frac{1}{T} \sum_t u_i(a_i, x_{-i}^t) \quad \forall a_i \in A_i \quad (1)$$

- Regret

$$R_i^T = \max_{a_i \in A_i} \sum_{t=1}^T u_i(a_i, x_{-i}^t) - \sum_{t=1}^T u_i(x_i^t, x_{-i}^t) \quad (2)$$

- Dividing Equation (2) by  $T$  and let  $T \rightarrow \infty$  yields Condition (1) since  $R_i^T/T$  tends to 0 by definition of no regret

Next lecture:

- Study a stronger regret notion called “**swap regret**” – it uses a stronger benchmark
- Show any game with no-swap-regret players will converge to a correlated equilibrium
- Prove that any no-regret algorithm can be converted to a no-swap-regret algorithm, with slightly worse regret guarantee

# Thank You

Haifeng Xu

University of Virginia

[hx4ad@virginia.edu](mailto:hx4ad@virginia.edu)