CS6501: Topics in Learning and Game Theory (Fall 2019)

Introduction to Mechanism Design

Instructor: Haifeng Xu

Outline

- Mechanism Design: Motivation and Examples
- Example Mechanisms for Single Item Allocation
- Example Mechanisms for Multiple Items Allocation

Mechanism Design: the Science of Rule Making

Mechanism Design (MD): designing a game by specifying its rules to induce a desired outcome among strategic participants

- So far, you are given the game and look to compute its equilibrium
  - For example, use no-regret learning dynamics or LPs
- In mechanism design, you design the game

Mechanism Design: the Science of Rule Making

Mechanism Design (MD): designing a game by specifying its rules to induce a desired outcome among strategic participants

- So far, you are given the game and look to compute its equilibrium
  - For example, use no-regret learning dynamics or LPs
- In mechanism design, you design the game
  - Specify game rules, player payoffs, allowable actions, etc.
  - Objective is to induce desirable outcome, e.g., incentivizing socially good or fair behaviors, maximizing revenue if selling goods
  - Typically, want the game to be easy to play
    - you don’t want it to be PPAD-hard for players to solve!

Importance of Rule Making: Tale 1

Determining HW deadline

- Will not answer Yes even you do

Importance of Rule Making: Tale 1

Determining HW deadline

- ➢ Now the reactions change
- ➢ Might answer Yes even you do not, but that comes also with risk

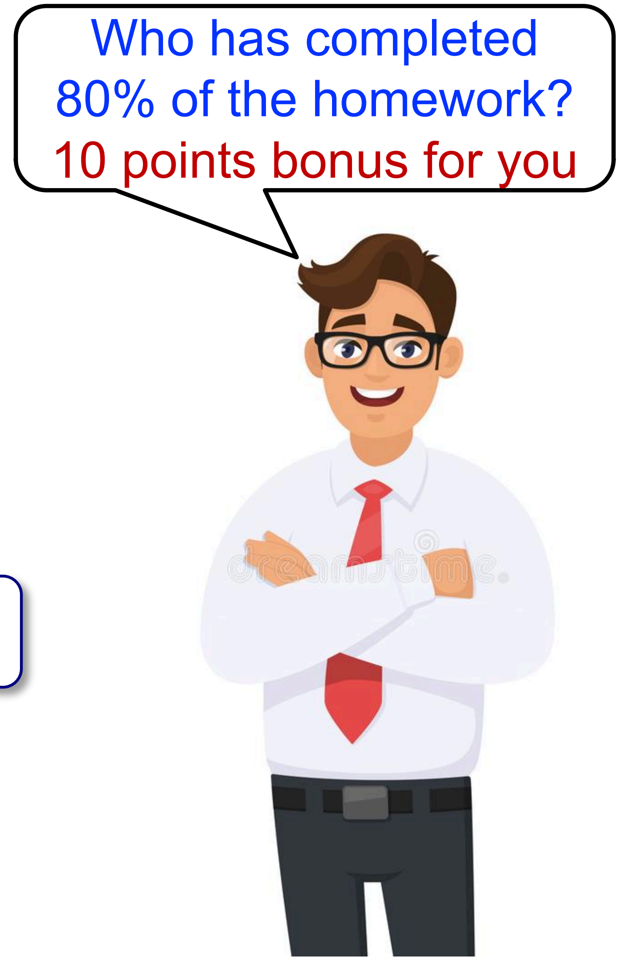

It is important to design the right rules!

Importance of Rule Making: Tale 2

A tale of horse racing

Importance of Rule Making: Tale 2

A tale of horse racing

- Two competitors; each has three horses of different levels: high, medium, low

Competitor 1

Competitor 2

high

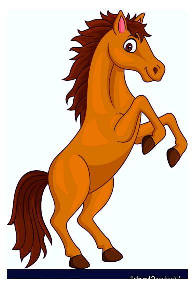

VS

medium

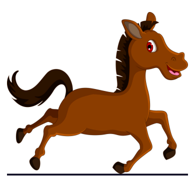

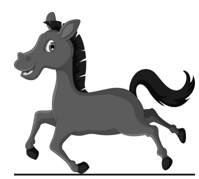

low

Importance of Rule Making: Tale 2

A tale of horse racing

- Two competitors; each has three horses of different levels: high, medium, low
- They need to compete at each horse level; whoever wins ≥ 2 times is the winner

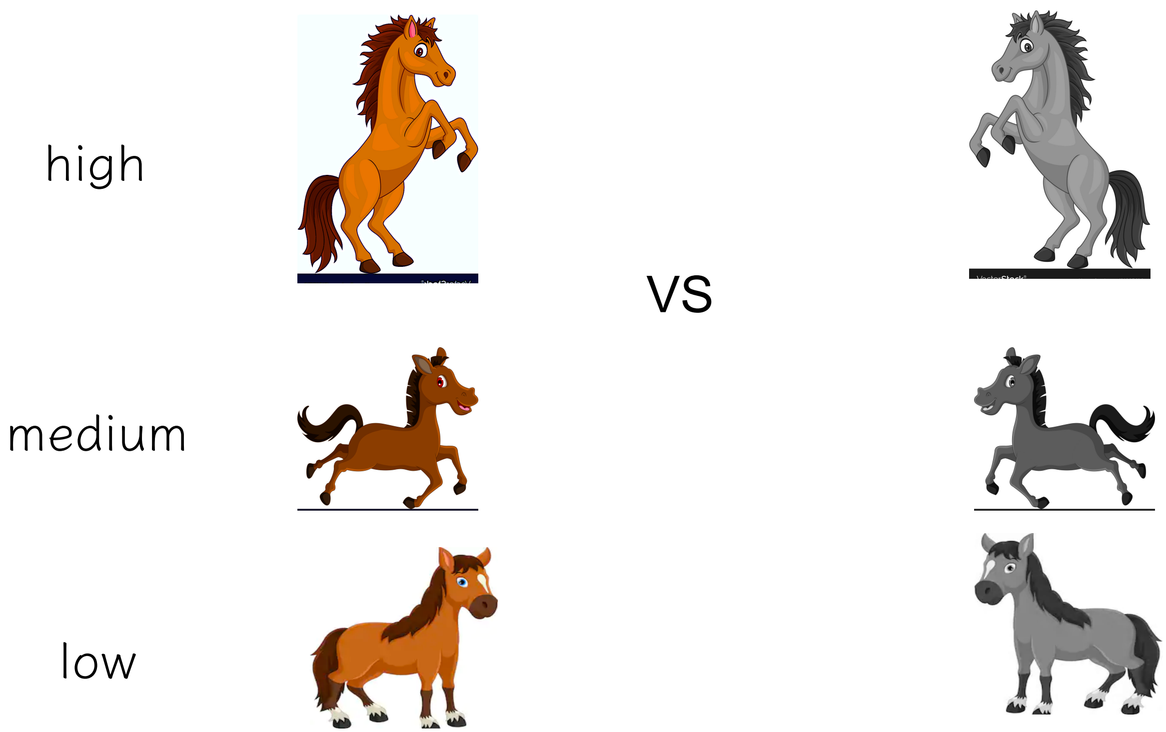

Importance of Rule Making: Tale 2

A tale of horse racing

- Two competitors; each has three horses of different levels: high, medium, low
- They need to compete at each horse level; whoever wins ≥ 2 times is the winner
  - Assume horses of different levels are indistinguishable but a horse at a higher level will always beat any horse at a lower level
- Can we truly determine the winner?
  - Both will look to use High horse against Medium and Medium against Low

Essentially no, winner will mainly depend on luck

Importance of Rule Making: Tale 2

A tale of horse racing

- Two competitors; each has three horses of different levels: high, medium, low
- What about the following rule?
  - They compete for 3 rounds
  - Winner of first round gains 3 points, winner of second round gains 2 points, and winner of the last round gains 1 point
  - Whoever gets the most points win

This is better – they will really compete at each level

Importance of Rule Making: Tale 3

- ➢ Selling products
  - Post a price
  - Customers only get to choose buy or not buy
- ➢ Why not the following mechanism?

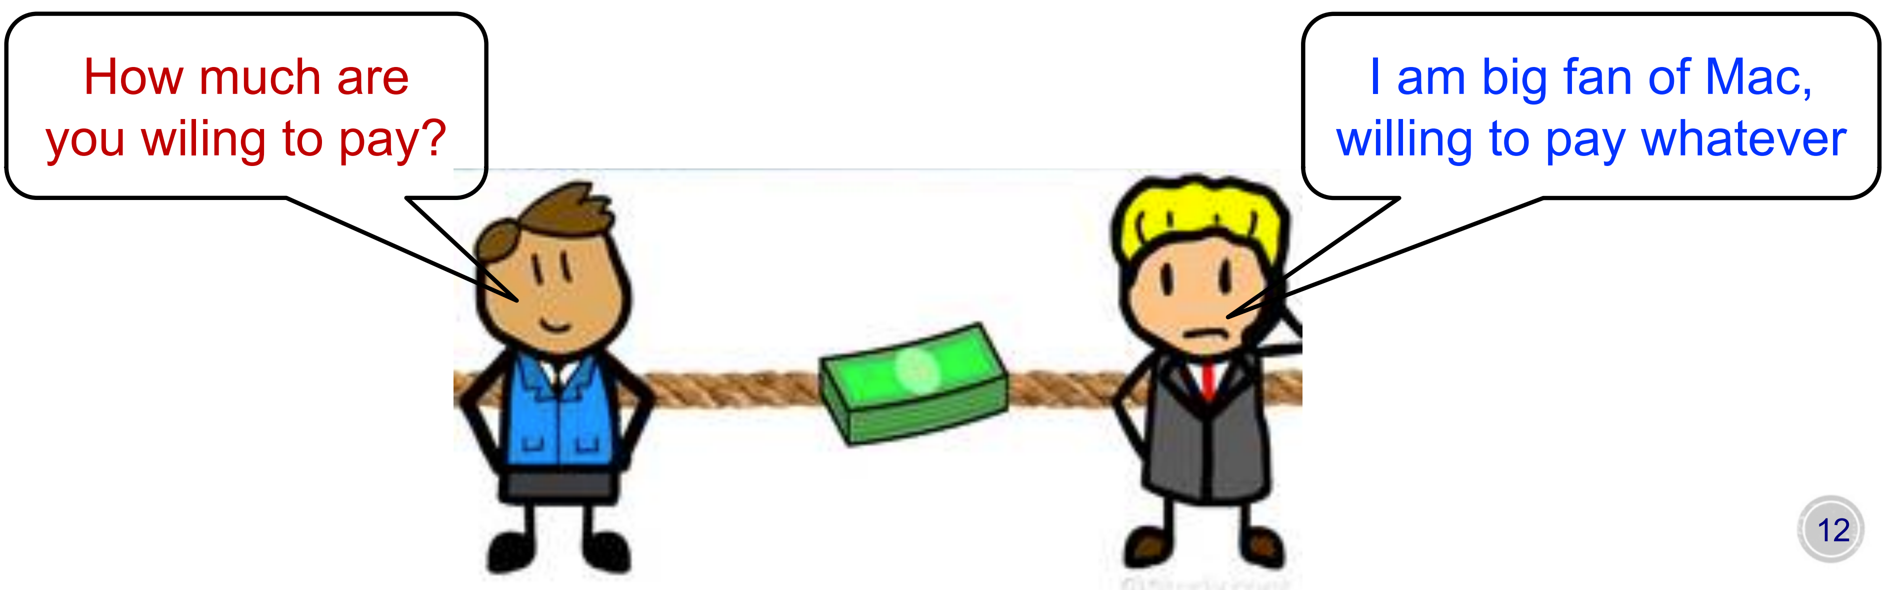

Importance of Rule Making: Tale 3

- ➢ Selling products
  - Post a price
  - Customers only get to choose buy or not buy
- ➢ Why not the following mechanism?
  - Customers will not be so honest

Ok, $2000 then

I am big fan of Mac, willing to pay whatever

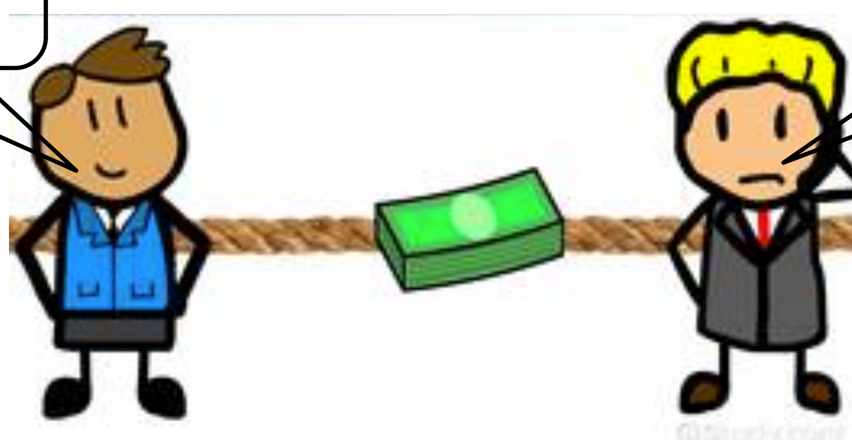

Importance of Rule Making: Tale 3

- Selling products
  - Post a price
  - Customers only get to choose buy or not buy
- Why not the following mechanism?
  - Customers will not be so honest
- Why not the following “bargaining” mechanism?

I have a price in mind, let me see whether your value is higher

My value is $2000

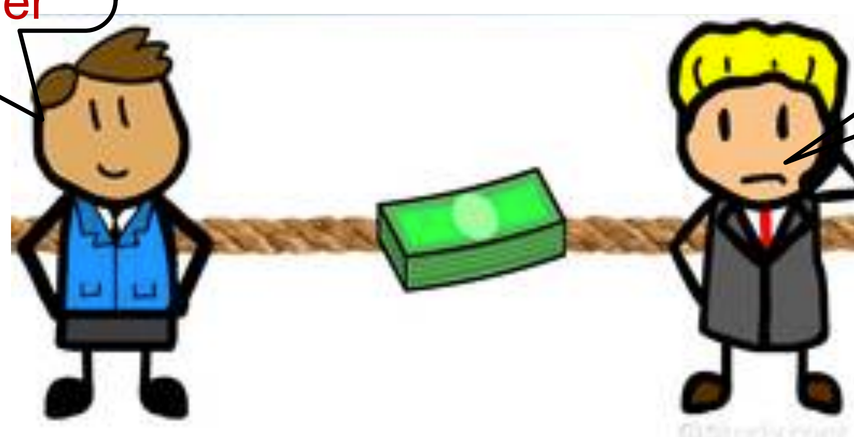

Importance of Rule Making: Tale 3

- Selling products
  - Post a price
  - Customers only get to choose buy or not buy
- Why not the following mechanism?
  - Customers will not be so honest
- Why not the following “bargaining” mechanism?

Ok, I will sell it to you for $1999

~~My value is $2000~~

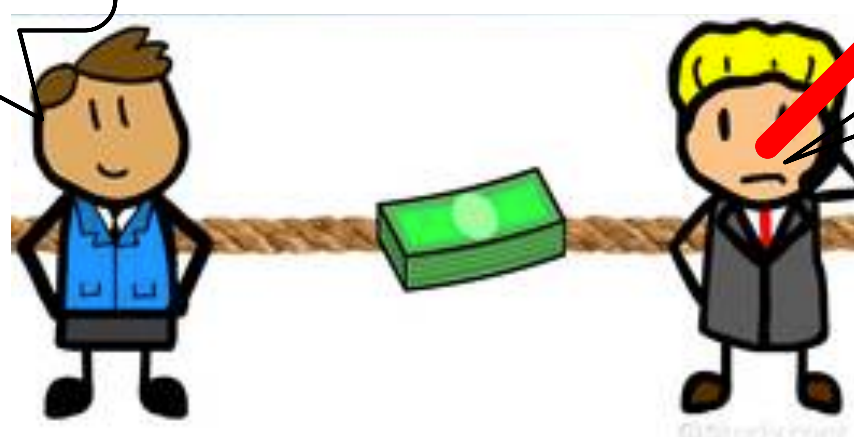

Importance of Rule Making: Tale 3

- ➢ Selling products
  - Post a price
  - Customers only get to choose buy or not buy
- ➢ Why not the following mechanism?
  - Customers will not be so honest
- ➢ Why not the following "bargaining" mechanism?

This is what's really going to happen…

I have a price in mind, let me see whether your value is higher

My value is $1000

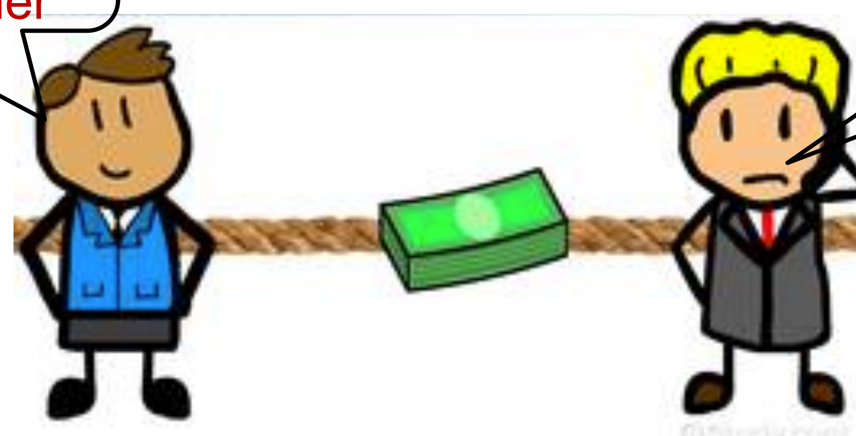

Importance of Rule Making: Tale 3

- Selling products
  - Post a price
  - Customers only get to choose buy or not buy
- Why not the following mechanism?
  - Customers will not be so honest
- Why not the following "bargaining" mechanism?

This is what's really going to happen…

No, no, won't sell

My value is $1000

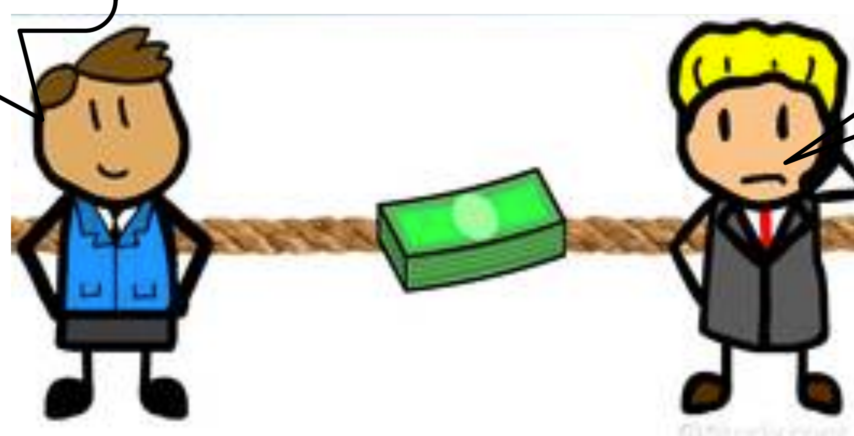

Importance of Rule Making: Tale 3

- Selling products
  - Post a price
  - Customers only get to choose buy or not buy
- Why not the following mechanism?
  - Customers will not be so honest
- Why not the following “bargaining” mechanism?

This is what’s really going to happen…

No, no, won’t sell

My value is $1001

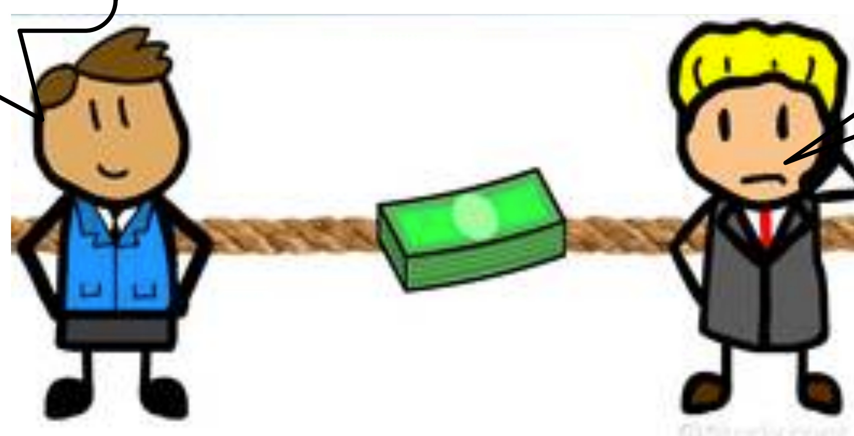

Importance of Rule Making: Tale 3

- Selling products
  - Post a price
  - Customers only get to choose buy or not buy
- Why not the following mechanism?
  - Customers will not be so honest
- Why not the following "bargaining" mechanism?

This is what's really going to happen…

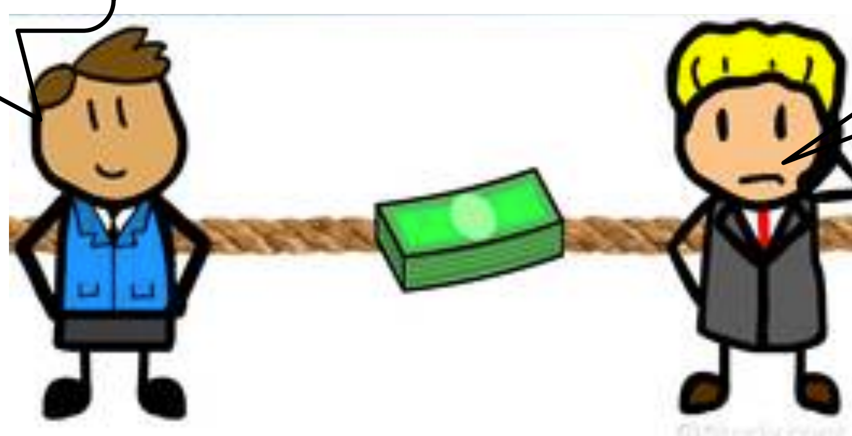

Importance of Rule Making: Tale 3

- Selling products
  - Post a price
  - Customers only get to choose buy or not buy
- Why not the following mechanism?
  - Customers will not be so honest
- Why not the following "bargaining" mechanism?

This is what's really going to happen…

No, no, won't sell

My value is ….

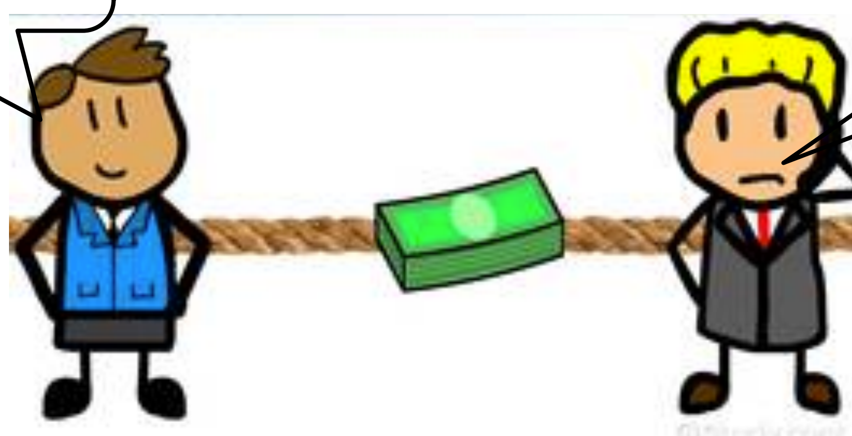

Importance of Rule Making: Tale 3

- Selling products
  - Post a price
  - Customers only get to choose buy or not buy
- Why not the following mechanism?
  - Customers will not be so honest
- Why not the following "bargaining" mechanism?

This is what's really going to happen…

Ok, I will sell it to you for $1900

My value is 1900

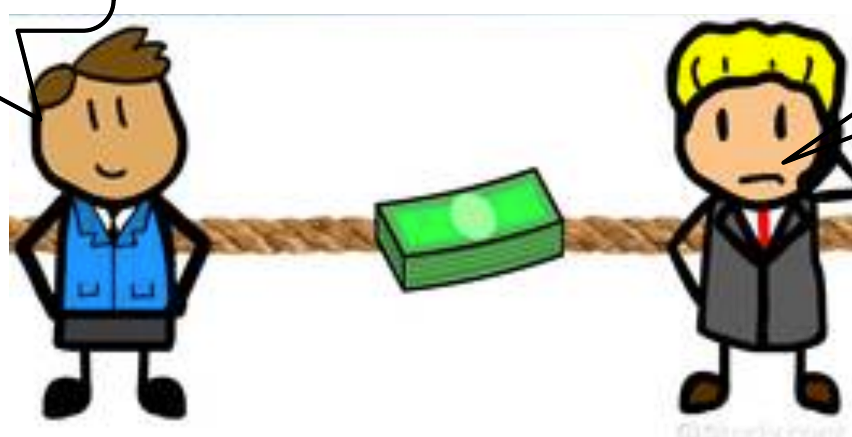

Importance of Rule Making: Tale 3

- Selling products
  - Post a price
  - Customers only get to choose buy or not buy
- Why not the following mechanism?
  - Customers will not be so honest
- Why not the following "bargaining" mechanism?
  - Too intricate buyer behaviors, interactions are too costly

Later, we will learn under mild assumptions, posting a price is optimal among all possible ways of selling to a buyer

Examples of Mechanism Design Problems

Example 1: Single-Item Allocation

v_1

v_2

- A single and indivisible item, n agents
- Agent i has a (private) value v_i about the item
- Outcome: choice of the winner of the item, and possibly payment from each agent
  - Note: payments do not have to involve, e.g., allocating temporary residence to homeless individuals
- Typical objectives: maximize revenue, maximize social welfare (i.e., allocate to the one who values the item most)
- Applications: selling items (e.g., e-bay), allocating scarce resources

Example 2: Multi-Item Allocation

- m items and n agents
- Agent i has (private) value $v_i(S)$ for any subset of items $S \subseteq [m]$
- Outcome: a partition of the items $[m]$ into $S_1, S_2, \ldots, S_n$ and agent i gets items in S_i
- Typical objectives: revenue, welfare, fairness
- Applications: rental room assignments, sell multiple products, dividing inheritance, etc.

Example 3: School Choice

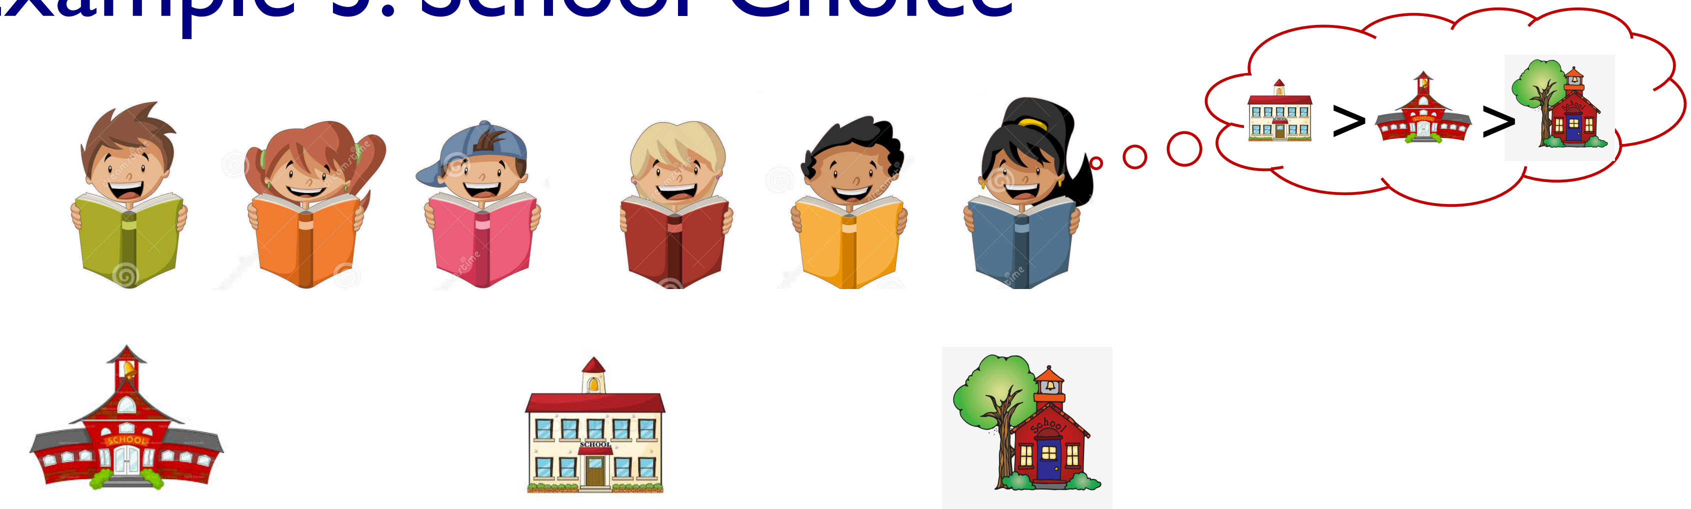

- n students, m schools
- Each student has a (private) preference over schools
  - Preference ≠ value function as in previous item allocation
- Similarly, each school also has a (private) preference over students
- Outcome: match each student to a school
- Objective: maximize “happiness” or “fairness”
- Applications: school choice, marriage or online dating, job matching, assigning web users to distributed Internet services, etc.

Example 4: Voting

- n voters, k candidates
- Each voter has a (private) preference over candidates
- Outcome: choice of a winning candidate
- Objective: maximize certain "social choice" function

Some Common Features

- Participants have private information (often called private types)
- Design objective typically depends on the private information
- Usually have to elicit such private information
- Participates are self-interested – they want to maximize their own utilities and may lie about their private information if helpful
  - Will be clear after we introduce mechanisms later

Outline

- Mechanism Design: Motivation and Examples
- Example Mechanisms for Single Item Allocation
- Example Mechanisms for Multiple Items Allocation

Single-Item Allocation

- A single and indivisible item, n agents
- Agent i has a (private) value v_i about the item
- Outcome: choice of the winner of the item, and possibly payment from each agent
- Typical objectives: maximize revenue, maximize social welfare

Benign Designer: Welfare Maximization

- Want to give the item to the agent who values it the most, i.e., $i^* = \arg\max_{i \in [n]} v_i$
  - But v_i is i's private information
  - The mechanism needs to elicit this information
  - Do not care about revenue
- Each agent is self-interested and wants to maximize his own utility $(v_i - p_i)$ where p_i is his payment (if any)

Benign Designer: Welfare Maximization

Q: what mechanism would work?

Trial 1: ask i to report his value b_i for all i; give the item to $i^* = \arg\max_i b_i$ (no payment)

- Use b_i because it may not equal v_i since agents may misreport
- Indeed, every one will report ∞

Can be proved that any mechanism without using payment cannot achieve the goal of welfare maximization.

Ok, need payment, so what is a natural mechanism with payment?

Benign Designer: Welfare Maximization

Q: what mechanism would work?

Trial 2: ask i to report his value b_i for all i; give the item to $i^* = \arg\max_i b_i$ and asks him to pay his own bid b_{i^*}

- This is called first-price auction
  - b_i called the "bid" and agents called the "bidders"
- Would agent report $b_i = v_i$?
  - They don't want → unnecessarily paying too much
  - They dare not report too small neither → may miss out on the item
  - Lead to very intricate and unpredictable agent behaviors
  - Winner does not necessarily have the highest v_i

Benign Designer: Welfare Maximization

Q: what mechanism would work?

Trial 3: ask i to report his value b_i for all i; give the item to $i^* = \arg\max_i b_i$ and asks him to pay the second highest bid $\max2_{i \in [n]} b_i$

- This is called second-price auction

Fact. Truthful bidding is a dominant strategy equilibrium in second-price auctions.

Intuition

- Fundamental reason: i'th payment does not depend on his own bid
  - i'th payment (if he wins) = highest bid among other bidders
  - So bid only affects whether i wins or not
  - Don't want to bid $b_i > v_i$ since that may make me pay more than v_i
  - Don't want to bid $b_i < v_i$ since whatever that bid wins, v_i also wins

Benign Designer: Welfare Maximization

Q: what mechanism would work?

Trial 3: ask i to report his value b_i for all i; give the item to $i^* = \arg\max_i b_i$ and asks him to pay the second highest bid $\max2_{i\in[n]} b_i$

- This is called second-price auction

Fact. Truthful bidding is a dominant strategy equilibrium in second-price auctions.

Formal proof:

- Fix a bidder i with true value v_i; let $b^* =$ highest bid among other bidders
- If $b^* < v_i$, any $b_i > b^*$ wins the item and pays b^*. So $b_i = v_i$ is also good
- If $b^*_{-i} \geq v_i$, i prefers losing. Bidding $b_i = v_i$ indeed will make him lose
- Though i does not know b^*, the reasoning above shows bidding $b_i = v_i$ is always optimal for whatever b^*

Benign Designer: Welfare Maximization

Q: what mechanism would work?

Trial 3: ask i to report his value b_i for all i; give the item to $i^* = \arg\max_i b_i$ and asks him to pay the second highest bid $\max2_{i\in[n]} b_i$

- This is called second-price auction

Fact. Truthful bidding is a dominant strategy equilibrium in second-price auctions.

- Thus truthful bidding are expected in second-price auctions
- So we will indeed give the item to the one with highest value
- This is the prototype of modern Ad Auctions used by Google, Microsoft, and many other ad exchange platforms
  - Reduces gaming behaviors in ad auctions

What About Revenue-Maximizing Designer?

- Studied much more in the literature
  - More motivated for designers with economic incentives
  - Welfare-maximization has been largely resolved
  - Revenue-maximization turns out to be much more difficult
  - Will also be our main focus in later lectures
- Without additional assumptions, cannot obtain any guarantee
  - Typically, need to assume prior knowledge about each bidder's value
  - Under natural assumptions, can be proved that optimal auction is roughly like a second-price auction, but with a "reserve price"
    - This should be surprising as there are really tons of ways to sell an item
    - This elegant auction format is optimal among all these ways
- Next, we show a simple example
  - Will see why second-price auction alone will not work

Example: Sell to Two Uniform Bidders

- Two bidders; for $i = 1,2$, $v_i \sim U([0,1])$ independently
- What is the expected revenue of second price auction?
  - Since bidders bid truthfully, revenue equals the smaller bidder value

$$Rev = \mathbb{E}_{v_1,v_2} \min(v_1, v_2) = 1/3$$

- Consider the following slight auction variant: highest bidder still wins, but pays $\max(second\ highest\ bid,\ 1/2)$
  - If both v_1, v_2 are less than $1/2$, keep the item with no sale
  - $1/2$ is called the "reserve price"
  - Truthful bidding is still a dominant strategy (the same proof)

Example: Sell to Two Uniform Bidders

- Two bidders; for $i = 1,2$, $v_i \sim U([0,1])$ independently
- What is the expected revenue of second price auction?
  - Since bidders bid truthfully, revenue equals the smaller bidder value

$$Rev = \mathbb{E}_{v_1,v_2} \min(v_1, v_2) = 1/3$$

- Consider the following slight auction variant: highest bidder still wins, but pays $\max(\textit{second highest bid},\ 1/2)$
- What is the expected revenue of this modified auction

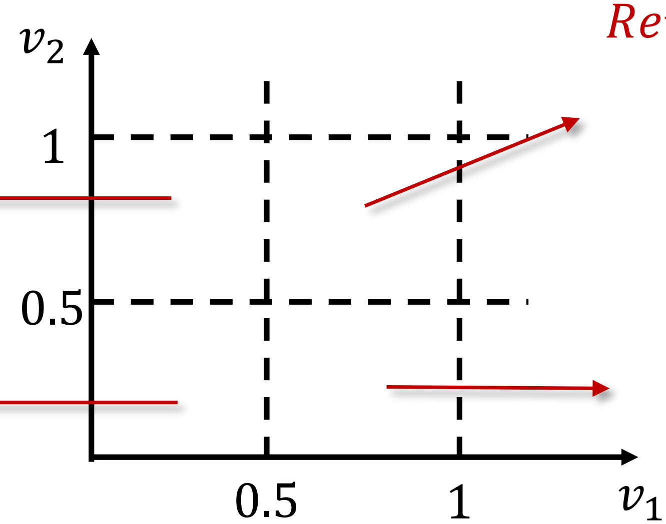

Sell to bidder 2, pay $\frac{1}{2}$, $Rev\ \frac{1}{8}$

0 revenue

$$Rev = \mathbb{E}_{v_1 \geq \frac{1}{2}, v_2 \geq \frac{1}{2}} \min(v_1, v_2)$$
$$= 1/6$$

Sell to bidder 1, pay $\frac{1}{2}$, $Rev\ \frac{1}{8}$

Example: Sell to Two Uniform Bidders

- Two bidders; for $i = 1,2$, $v_i \sim U([0,1])$ independently
- What is the expected revenue of second price auction?
  - Since bidders bid truthfully, revenue equals the smaller bidder value

$$Rev = \mathbb{E}_{v_1,v_2} \min(v_1, v_2) = 1/3$$

- Consider the following slight auction variant: highest bidder still wins, but pays $\max(second\ highest\ bid,\ 1/2)$
- What is the expected revenue of this modified auction
  - Total revenue is $\frac{1}{8} + \frac{1}{8} + \frac{1}{6} = \frac{5}{12}$, which turns out to be optimal revenue
  - Second price auction is not optimal because it charges too little when $v_1 > 1/2 > v_2$
  - ½ here is not arbitrary → it equals $\arg \max_{x \in [0,1]} x(1 - F(x))$

Outline

- Mechanism Design: Motivation and Examples
- Example Mechanisms for Single Item Allocation
- Example Mechanisms for Multiple Items Allocation

Multi-Item Allocation

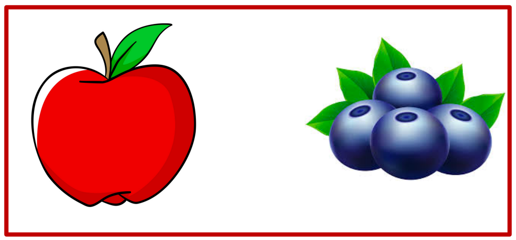

S_1

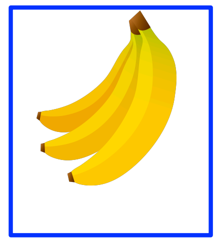

S_2

$v_1(S)$

$v_2(S)$

- m items and n agents
- Agent i has (private) value $v_i(S)$ for any subset of items $S \subseteq [m]$
- Outcome: a partition of the items $[m]$ into $S_1, S_2, \dots, S_n$ and agent i gets items in S_i
- Typical objectives: revenue, welfare, fairness
  - Revenue-maximizing is extremely challenging – huge amount of research, still a major open question in economics and CS
  - A lot of study on fair allocation as well – challenging in general
  - Welfare maximization can be solved via an elegant generalization of second-price auction

Multi-Item Allocation: Welfare Maximization

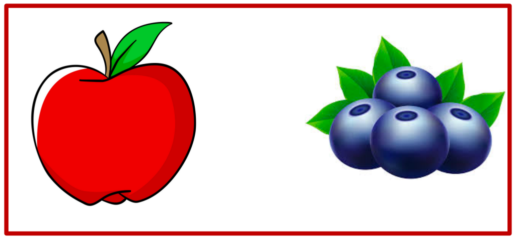

S_1^*

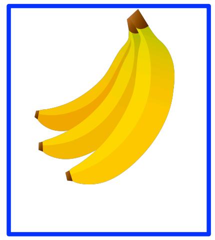

S_2^*

$v_1(S)$

$v_2(S)$

- The Vickrey-Clarke-Groves (VCG) mechanism
  1. Ask each bidder to report their value function $b_i(S)$
  2. Compute optimal allocation $(S_1^*, \cdots, S_n^*) = \arg\max_{(S_1,\cdots,S_n)} \sum_{i=1}^n b_i(S_i)$
  3. Allocate S_i^* to bidder i, charge i the following amount

$$p_i = \underbrace{\left[\max_{S_{-i}} \sum_{j \neq i} b_j(S_j)\right]}_{\text{Maximum welfare if } i \text{ did not participate}} - \underbrace{\sum_{j \neq i} b_j(S_j^*)}_{\text{Other's welfare when } i \text{ participates}}$$

p_i = how much i "hurts" all the others' welfare due to his participation

Multi-Item Allocation: Welfare Maximization

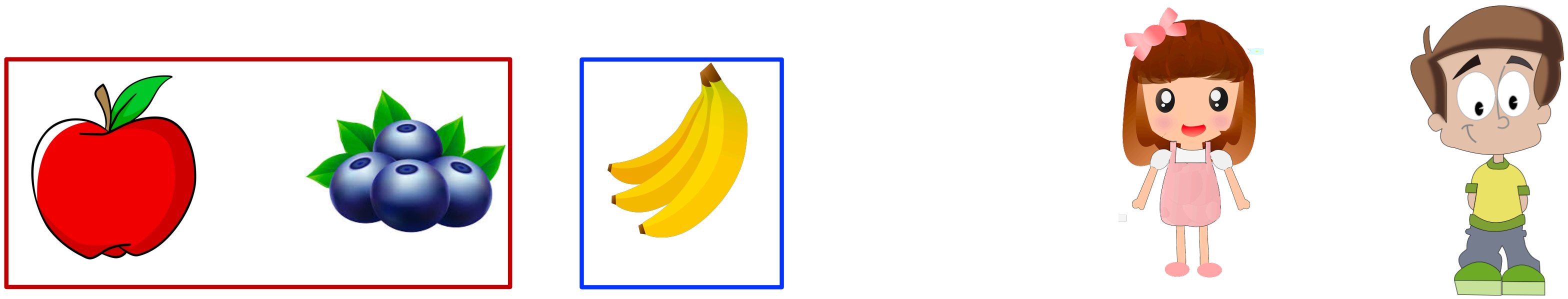

S_1 S_2 $v_1(S)$ $v_2(S)$

- The Vickrey-Clarke-Groves (VCG) mechanism
  1. Ask each bidder to report their value function $b_i(S)$
  2. Compute optimal allocation $(S_1^*, \cdots, S_n^*) = \arg \max_{(S_1, \cdots, S_n)} \sum_{i=1}^n b_i(S_i)$
  3. Allocate S_i^* to bidder i, charge i the following amount

$$p_i = \left[\max_{S_{-i}} \sum_{j \neq i} b_j(S_j) \right] - \sum_{j \neq i} b_j(S_j^*)$$

Q: what is p_i if there is only a single item for sale?

1. The item will be allocated to largest $b_i(\text{item})$
2. Winner pays the second highest bid; others pay 0
3. Degenerate to a second price auction

Multi-Item Allocation: Welfare Maximization

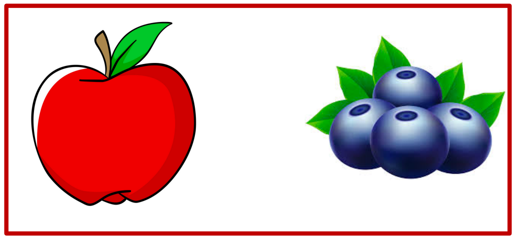

S_1

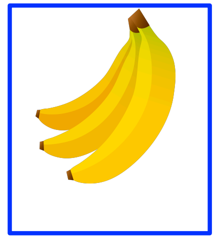

S_2

$v_1(S)$

$v_2(S)$

- The Vickrey-Clarke-Groves (VCG) mechanism
  1. Ask each bidder to report their value function $b_i(S)$
  2. Compute optimal allocation $(S_1^*, \cdots, S_n^*) = \arg\max_{(S_1,\cdots,S_n)} \sum_{i=1}^{n} b_i(S_i)$
  3. Allocate S_i^* to bidder i, charge i the following amount

$$p_i = \left[\max_{S_{-i}} \sum_{j \neq i} b_j(S_j)\right] - \sum_{j \neq i} b_j(S_j^*)$$

Fact. Truthful bidding is a dominant strategy equilibrium in VCG.

- So it does maximize welfare at equilibrium
- Proof: HW exercise

Thank You

Haifeng Xu

University of Virginia

hx4ad@virginia.edu