Haifeng Xu

Assistant Professor of Computer Science, University of Chicago
Crerar 260, 5730 S Ellis Ave, Chicago, IL 60637.

✉ haifengxu@uchicago.edu   % www.haifeng-xu.com   ☞ Google Scholar

RESEARCH INTERESTS

Computational Economics, Machine Learning, Artificial Intelligence, Multi-Agent Systems, the Design and Analysis of Algorithms

PROFESSIONAL EXPERIENCE

Assistant Professor   07/2022 – present
Department of Computer Science
University of Chicago, USA

Alan Batson Assistant Professor   08/2019 – 06/2022
Department of Computer Science
University of Virginia, USA

Visiting Research Scientist (one day per week)   08/2021 – 02/2022
Market Algorithm Research Group, Google, USA

Postdoctoral Fellow   08/2018 – 08/2019
Center for Research on Computation and Society (CRCS), *Harvard University, USA*
Hosts: Yiling Chen and David C. Parkes

Lecturer in Computer Science, Harvard University   09/2018 – 12/2018

Research Intern, Google Research   06/2016 – 08/2016

Visiting Scholar, Simons Institute for the Theory of Computing   10/2015 – 12/2015
  Program: Economics and Computation

Research Intern, Yahoo! Lab   06/2015 – 08/2015

Research Intern, Microsoft Research   07/2011 – 06/2012

EDUCATION

Ph.D. in Computer Science   08/2013 – 07/2018
University of Southern California, USA
Advisors: Milind Tambe and Shaddin Dughmi
Thesis: Information as a Double-Edged Sword in Strategic Interactions.
ACM SIGecom Doctoral Dissertation Award (honorable mention)
IFAAMAS Victor Lesser Distinguished Dissertation Award (runner-up)

MMath in Computational Mathematics   08/2012 – 08/2013
University of Waterloo, Canada

B.Sc. (honours) in Mathematics   08/2008 – 07/2012
School of Gifted Young, University of Science & Technology of China
I was a member of the HUA Loo-Keng Elite Program in Mathematics

HONORS & AWARDS

○ Google Faculty Research Award 2020
  One of the six recipients in *Algorithms and Optimization* within North America

○ ACM SIGecom Dissertation Award (honorable mention) 2019

○ IFAAMAS Victor Lesser Distinguished Dissertation Award (runner-up) 2019

○ Best Application System Demo Award, AAMAS 2019. 2019

○ Google Ph.D. Fellowship 2017
  One of the three recipients worldwide in the category *Algorithms, Optimization and Markets*

○ CAMS Prize for Excellence in Research, USC Center for Applied Mathematical Sciences 2017
  Awarded annually to two graduate students across USC

○ Best Paper Award, AAMAS Workshop on Security and Multi-agent Systems (SecMas) 2016

○ Best Student Paper Award, AAMAS 2016 2016

○ Shing-Tung Yau Chinese College Student Mathematics Contests 2011
  A prestigious national contest organized by the famous mathematician Shing-Tung Yau
  – Silver Medal in Applied Mathematics (Top 4 in the country)
  – Bronze Team Medal (Top 4 teams in the country)

PUBLICATIONS

*α-β indicates alphabetical author order, which is typical for theoretical CS venues

Papers Under Review

[U1]. You Zu, Krishnamurthy Iyer and **Haifeng Xu**. Learning to Persuade on the Fly: Robustness Against Ignorance. *R&R to **Operations Research***.

Journal Papers

[J5]. (α-β) Ravi Sundaram, Anil Vullikanti, **Haifeng Xu** and Fan Yao. PAC-Learning for Strategic Classification. ***Journal of Machine Learning Research***, conditional acceptance with minor revision.

[J4]. (α-β) Yakov Babichenko, Inbal Talgam-Cohen, **Haifeng Xu** and Konstantin Zabarnyi. Regret-Minimizing Bayesian Persuasion. ***Games and Economic Behavior***, 2022.

[J3]. **Haifeng Xu**, Rasha Kashef, Hans De Sterck and Geoffrey Sanders. Efficient Algebraic Multigrid Methods for Multilevel Overlapping Co-Clustering of User-Item Relationships. ***INFORMS Journal on Computing***, 2022.

[J2]. (α-β) Shaddin Dughmi, **Haifeng Xu.** Algorithmic Bayesian Persuasion. ***SIAM Journal on Computing***, *page STOC16-68, 2019.*

[J1]. Amulya Yadav, Hau Chan, Albert Jiang, **Haifeng Xu**, Eric Rice and Milind Tambe. Using Social Networks to Raise HIV Awareness Among Homeless Youth. ***IBM Journal of Research and Development***: *Volume 61, Issue 6(4), pp. 1-10, 2017.*

Refereed Conference Papers

[C60]. (α-β) Bolin Ding, Yiding Feng, Chien-Ju Ho, Wei Tang and **Haifeng Xu**. Competitive Information Design for Pandora's Box. *Proceedings of the 34th ACM-SIAM Symposium on Discrete Algorithms*, (**SODA 2023**).

[C59]. Jibang Wu, Weiran Shen, Fei Fang and **Haifeng Xu**. Inverse Game Theory for Stackelberg Games: the Blessing of Bounded Rationality. *Proceedings of the 36th Conference on Neural Information Processing Systems* (**NeurIPS'22**).

[C58]. Ashwinkumar Badanidiyuru, Zhe Feng, Tianxi Li and **Haifeng Xu**. Incrementality Bidding via Reinforcement Learning under Mixed and Delayed Rewards. *Proceedings of the 36th Conference on Neural Information Processing Systems* (**NeurIPS'22**).

[C57]. Stephanie Schoch, **Haifeng Xu** and Yangfeng Ji. CS-Shapley: Class-wise Shapley Values for Data Valuation in Classification. *Proceedings of the 36th Conference on Neural Information Processing Systems* (**NeurIPS'22**).

[C56]. (α-β) Jibang Wu, Fan Yao and **Haifeng Xu**. Multi-Agent Learning for Iterative Dominance Elimination: Formal Barriers and New Algorithms. *Proceedings of the 35th Annual Conference on Learning Theory*, (**COLT 2022**).

[C55]. Jibang Wu, Zixuan Zhang, Zhe Feng, Zhaoran Wang, Zhuoran Yang, Michael I. Jordan and **Haifeng Xu**. Sequential Information Design: Markov Persuasion Process and Its Efficient Reinforcement Learning. *Proceedings of the 23th ACM Conference on Economics and Computation*, (**EC 2022**).

[C54]. Chenghan Zhou, Thanh H. Nguyen and **Haifeng Xu**. Algorithmic Information Design in Multi-Player Games: Possibility and Limits in Singleton Congestion. *Proceedings of the 23th ACM Conference on Economics and Computation*, (**EC 2022**).

[C53]. (α-β) Yiding Feng, Wei Tang and **Haifeng Xu**. Online Bayesian Recommendation with No Regret. *Proceedings of the 23th ACM Conference on Economics and Computation*, (**EC 2022**).

[C52]. Fan Yao, Chuanhao Li, Denis Nekipelov, Hongning Wang and **Haifeng Xu**. Learning from a Learning User for Optimal Recommendations. *Proceedings of the 39th International Conference on Machine Learning*, (**ICML 2022**).

[C51]. (α-β) Junjie Chen, Minming Li and **Haifeng Xu**. Selling Data To a Machine Learner: Pricing via Costly Signaling. *Proceedings of the 39th International Conference on Machine Learning*, (**ICML 2022**).

[C50]. Huazheng Wang, **Haifeng Xu** and Hongning Wang. When Are Linear Stochastic Bandits Attackable? *Proceedings of the 39th International Conference on Machine Learning*, (**ICML 2022**).

[C49]. Chenghan Zhou, Andrew Spivey, **Haifeng Xu** and Thanh Hong Nguyen. Information Design for Multiple Independent and Self-Interested Defenders: Work Less, Pay Off More. *Proceedings of the 38th Conference on Uncertainty in Artificial Intelligence*, (**UAI 2022**).

[C48]. (α-β for first two authors) Anshuka Rangi, **Haifeng Xu**, Long Tran-Thanh and Massimo Franceschetti. Understanding the Limits of Poisoning Attacks in Episodic Reinforcement Learning. *Proceedings of the 31st International Joint Conference on Artificial Intelligence* (**IJCAI 2022**), long oral, 3.75%

[C47]. (α-β) Quinlan Dawkins, Minbiao Han and **Haifeng Xu**. First-order Convex Fitting and Its Application to Economics and Optimization. *Proceedings of the 36nd AAAI Conference on Artificial Intelligence* (**AAAI 2022**, acceptance rate 15%), long oral

[C46]. Fan Yao, Chuanhao Li, Denis Nekipelov, Hongning Wang and **Haifeng Xu**. Learning the Optimal Recommendation from Explorative Users. *Proceedings of the 36nd AAAI Conference on Artificial Intelligence* (**AAAI 2022**, acceptance rate 15%)

[C45]. **Haifeng Xu** and Ruggiero Cavallo. The Strange Role of Information Asymmetry in Auctions — Does More Accurate Value Estimation Benefit A Bidder? *Proceedings of the 36nd AAAI Conference on Artificial Intelligence* (**AAAI 2022**, acceptance rate 15%), long oral

[C44]. (α-β) Anshuka Rangi, Long Tran-Thanh, **Haifeng Xu** and Massimo Franceschetti. Saving Stochastic Bandits from Poisoning Attacks via Limited Data Verification. *Proceedings of the 36nd AAAI Conference on Artificial Intelligence* (**AAAI 2022**, acceptance rate 15%), long oral

[C43]. Thanh Nguyen and **Haifeng Xu**. When Can the Defender Effectively Deceive Attackers in Security Games? *Proceedings of the 36nd AAAI Conference on Artificial Intelligence* (**AAAI 2022**, acceptance rate 15%).

[C42]. (α-β) Yakov Babichenko, Inbal Talgam-Cohen, **Haifeng Xu** and Konstantin Zabarnyi. Multi-Channel Bayesian Persuasion. *Proceedings of the Innovations in Theoretical Computer Science* (**ITCS'22**).

[C41]. (α-β) Quinlan Dawkins, Minbiao Han and **Haifeng Xu**. The Limits of Optimal Pricing in the Dark. *Proceedings of the 35th Conference on Neural Information Processing Systems* (**NeurIPS'21**).

[C40]. Sijun Tan, Jibang Wu, Xiaohui Bei and **Haifeng Xu**. Least Square Calibration for Peer Reviews. *Proceedings of the 35th Conference on Neural Information Processing Systems* (**NeurIPS'21**).

[C39]. Chengshuai Shi, **Haifeng Xu**, Wei Xiong, Cong Shen. (Almost) Free Incentivized Exploration from Decentralized Learning Agents. *Proceedings of the 35th Conference on Neural Information Processing Systems* (**NeurIPS'21**).

[C38]. (α-β) Ravi Sundaram, Anil Vullikanti, **Haifeng Xu** and Fan Yao. PAC-Learning for Strategic Classification. *Proceedings of the 38th International Conference on Machine Learning* (**ICML 2021**), long oral, 3%.

[C37]. Quinlan Dawkins, Tianxi Li and **Haifeng Xu**. Diffusion Source Identification on Networks with Statistical Confidence. *Proceedings of the 38th International Conference on Machine Learning* (**ICML 2021**).

[C36]. Shuze Liu, Weiran Shen and **Haifeng Xu**. Optimal Pricing of Information. *Proceedings of the 22th ACM Conference on Economics and Computation* (**EC'21**).

[C35]. (α-β) Yakov Babichenko, Inbal Talgam-Cohen, **Haifeng Xu** and Konstantin Zabarnyi. Regret-Minimizing Bayesian Persuasion. *Proceedings of the 22th ACM Conference on Economics and Computation* (**EC'21**).

[C34]. You Zu, Krishnamurthy Iyer and **Haifeng Xu**. Learning to Persuade on the Fly: Robustness Against Ignorance. *Proceedings of the 22th ACM Conference on Economics and Computation* (**EC'21**).

[C33]. James Nachbar and Haifeng Xu. The Power of Signaling and its Intrinsic Connection to the Price of Anarchy. *Proceedings of the Third International Conference on Distributed Artificial Intelligence*, (**DAI'21**).

[C32]. Aditya Mate, Jackson A. Killian, **Haifeng Xu**, Andrew Perrault, Milind Tambe. Collapsing Bandits and Their Application to Public Health Interventions. *Proceedings of the 34th Conference on Neural Information Processing Systems* (**NeurIPS 2020**).

[C31]. (α-β) Zhe Feng, David Parkes, **Haifeng Xu**. The Intrinsic Robustness of Stochastic Bandits to Strategic Manipulation. *Proceedings of the 37th International Conference on Machine Learning* (**ICML 2020**).

[C30]. **Haifeng Xu**. On the Tractability of Public Persuasion with No Externalities. *Proceedings of the 31st ACM-SIAM Symposium on Discrete Algorithms* (**SODA 2020**).

[C29]. (α-β) Yiling Chen, **Haifeng Xu**, Shuran Zheng. Selling Information through Consulting. *Proceedings of the 31st ACM-SIAM Symposium on Discrete Algorithms* (**SODA 2020**).

[C28]. Sarah Keren, **Haifeng Xu**, David Parkes, Kofi Kwapong and Barbara Grosz. Information Shaping for Enhanced Goal Recognition of Partially-Informed Agents. *Proceedings of the 34nd AAAI Conference on Artificial Intelligence* (**AAAI 2020**), oral.

[C27]. Elizabeth Bondi, Hoon Oh, **Haifeng Xu**, Fei Fang, Bistra Dilkina, Milind Tambe. To Signal or Not To Signal: Exploiting Uncertain Real-Time Information in Signaling Games for Security and Sustainability. *Proceedings of the 34nd AAAI Conference on Artificial Intelligence* (**AAAI 2020**), oral.

[C26]. Chao Yan, **Haifeng Xu**, Yevgeniy Vorobeychik, Bo Li, Daniel Fabbri, Bradley Malin. To Warn or Not to Warn: Online Signaling in Audit Games. *Proceedings of the 36th IEEE International Conference on Data Engineering* (**ICDE 2020**), oral.

[C25]. Thanh Nguyen, Andrew Butler and **Haifeng Xu**. Tackling Imitative Attacker Deception in Repeated Bayesian Stackelberg Security Games. *Proceedings of the 24th European Conference on Artificial Intelligence* (**ECAI 2020**).

[C24]. Jiarui Gan, **Haifeng Xu**, Qingyu Guo, Long Tran-Thanh, Zinovi Rabinovich and Michael Wooldridge. Imitative Follower Deception in Stackelberg Games. *Proceedings of the 20th ACM Conference on Economics and Computation* (**EC'19**).

[C23]. (α-β) Thanh H. Nguyen and **Haifeng Xu**. Imitative Attacker Deception in Stackelberg Security Games. *Proceedings of the 28th International Joint Conference on Artificial Intelligence* (**IJCAI'19**).

[C22]. (α-β) Jerry Anunrojwong, Yiling Chen, Bo Waggoner and **Haifeng Xu**. Computing Equilibria of Prediction Markets via Persuasion. *Proceedings of the 15th Conference on Web and Internet Economics* (**WINE'19**).

[C21]. Omkar Thakoor, Milind Tambe, Phebe Vayanos, **Haifeng Xu**, Christopher Kiekintveld, and Fei Fang. Cyber Camouflage Games for Strategic Deception. *Proc. 10th Conference on Decision and Game Theory for Security* (**GameSec 2019**).

[C20]. **Haifeng Xu**, Kai Wang, Phebe Vayanos, Milind Tambe. Strategic Coordination of Human Patrollers and Mobile Sensors with Signaling for Security Games. *Proceedings of the 32th AAAI Conference on Artificial Intelligence* (**AAAI'18**), oral.

[C19]. (α-β) Ashwinkumar Badanidiyuru, Kshipra Bhawalkar, **Haifeng Xu**. Targeting and Signaling in Ad Auctions. *ACM-SIAM Symposium on Discrete Algorithms* (**SODA'18**).

[C18]. **Haifeng Xu**, Shaddin Dughmi, Milind Tambe, Venil Loyd Noronha. Mitigating the Curse of Correlation in Security Games by Entropy Maximization. *Proceedings of the 17th International Conference on Autonomous Agents and Multiagent Systems* (**AAMAS'18**, short paper).

[C17]. Aaron Schlenker, Omkar Thakoor, **Haifeng Xu**, Fei Fang, Milind Tambe, Long Tran-Thanh, Phebe Vayanos, Yevgeniy Vorobeychik. Deceiving Cyber Adversaries: A Game Theoretic Approach. *Proceedings of the 17th International Conference on Autonomous Agents and Multiagent Systems* (**AAMAS'18**).

[C16]. (α-β) Shaddin Dughmi, **Haifeng Xu**. Algorithmic Persuasion with No Externalities. *Proceedings of the 18th ACM Conference on Economics and Computation* (**EC'17**).

[C15]. **Haifeng Xu***, Benjamin Ford*, Fei Fang, Bistra Dilkina, Andrew Plumptre, Milind Tambe, Margaret Driciru, Fred Wanyama, Aggrey Rwetsiba, Mustapha Nsubaga and Joshua Mabonga. Optimal Patrol Planning for Green Security Games with Black-Box Attackers. *Proceedings of the 8th Conference on Decision and Game Theory for Security* (**GameSec'17**). (*Equal Contributions)

[C14]. Aaron Schlenker, **Haifeng Xu**, Mina Guirguis, Christopher Kiekintveld, Arunesh Sinha, Milind Tambe, Solomon Sonya, Darryl Balderas, Noah Dunstatter. Don't Bury your Head in Warnings: A Game-Theoretic Approach for Intelligent Allocation of Cyber-security Alerts. *Proceedings of the 26th International Joint Conference on Artificial Intelligence* (**IJCAI'17**). Highlighted in the press release opening the IJCAI'17 conference.

[C13]. (α-β) Shaddin Dughmi, **Haifeng Xu**. Algorithmic Bayesian Persuasion. *Proceedings of the 48th ACM Symposium on Theory of Computing* (**STOC'16**). Invited to SICOMP Special Issue for STOC 2016.

[C12]. **Haifeng Xu.** The Mysteries of Security Games: Equilibrium Computation Becomes Combinatorial Algorithm Design. *Proceedings of the 17th ACM Conference on Economics and Computation* (**EC'16**). Best paper award at the AAMAS-16 workshop on Security and Multi-agent Systems.

[C11]. **Haifeng Xu**, Rupert Freeman, Vincent Conitzer, Shaddin Dughmi, Milind Tambe. Signaling in Bayesian Stackelberg Games. *Proceedings of the 15th International Conference on Autonomous Agents and Multiagent Systems* (**AAMAS'16**).

[C10]. **Haifeng Xu***, Long Tran Thanh*, Nick Jennings. Playing Repeated Security Games with No Prior Knowledge. *Proceedings of the 15th International Conference on Autonomous Agents and Multiagent Systems* (**AAMAS'16**). (*Equal Contributions)

[C9]. Amulya Yadav, Hau Chan, Albert Jiang, **Haifeng Xu**, Eric Rice, Milind Tambe. Using Social Networks to Aid Homeless Shelters: Dynamic Influence Maximization Under Uncertainty. *Proceedings of the 15th International Conference on Autonomous Agents and Multiagent Systems* (**AAMAS'16**). Best student paper award.

[C8]. **Haifeng Xu**, Albert X. Jiang, Arunesh, Sinha, Zinovi Rabinovich, Shaddin Dughmi, Milind Tambe. Security Games with Information Leakage: Modeling and Computation. *Proceedings of the 24th International Joint Conference on Artificial Intelligence* (**IJCAI'15**).

[C7]. Yue Yin, **Haifeng Xu**, Jiarui Gan, Bo An, Albert X. Jiang. Computing Optimal Mixed Strategies for Security Games With Dynamic Payoffs. *Proceedings of the 24th International Joint Conference on Artificial Intelligence* (**IJCAI'15**).

[C6]. Zinovi Rabinovich, Albert X. Jiang, Manish Jain, **Haifeng Xu**. Information Disclosure as a Means of Security. *Proceedings of the 14th International Conference on Autonomous Agents and Multiagent Systems* (**AAMAS'15**).

[C5]. **Haifeng Xu**, Zinovi Rabinovich, Shaddin Dughmi, Milind Tambe. Exploring Information Asymmetry in Two-Stage Security Games. *Proceedings of the 29th AAAI Conference on Artificial Intelligence* (**AAAI'15**).

[C4]. **Haifeng Xu**, Fei Fang, Albert X. Jiang, Vincent Conitzer, Shaddin Dughmi, Milind Tambe. Solving Zero-Sum Security Games in Discretized Spatio-Temporal Domains. *Proceedings of the 28th AAAI Conference on Artificial Intelligence* (**AAAI'14**).

[C3]. Leandro Marcolino, **Haifeng Xu**, Albert X. Jiang, Milind Tambe, Emma Bowring. Give a Hard Problem to a Diverse Team: Exploring Large Action Spaces. *Proceedings of the 28th AAAI Conference on Artificial Intelligence* (**AAAI'14**).

[C2]. **Haifeng Xu**, Kate Larson. Improving the Efficiency of Crowdsourcing Contests. *Proceedings of the 13th International Conference on Autonomous Agents and Multiagent Systems* (**AAMAS'14**).

[C1]. **Haifeng Xu**, Bin Gao, Diyi Yang, Tieyan Liu. Predicting Advertiser Bidding Behaviors in Sponsored Search by Rationality Modeling. *Proceedings of the 22nd International Conference on World Wide Web* (**WWW'13**).

SPONSORED RESEARCH AND PROGRAMS

*Total: 1005K; My Share: 970K

Years	Description	Total	My Share
2023-2026	ARO W911NF2310030: *Efficient Deception-Aware Learning of Game-Theoretic Solutions for Adversarial Domains*, sole PI: Xu	$416,385	$416,385
2021-2024	NSF Award 2303372: *AF: Algorithmic Persuasion: Re-creating the Success of Mechanism Design*, sole PI: Xu	$453,382	$453,382
2021-2022	Google AI for Social Good Award: *Combating Poaching through Community Influence*. PI: Xu, co-PI: World Wildlife Fund	$30,000	$10,000
2020-2021	Google Faculty Research Award: *Boosting Ad Auctions through Better Information Design*. sole PI: Xu.	$55,000	$55,000

| 2020-2021 | Global Infectious Diseases Institute: *Impact of Strategic Behavior on Inequities in Health Outcomes during COVID-19*. PI: Xu, co-PI: Deborah Hellman, Achla Marathe, Anil Vullikanti. | $50,000 | $35,000 |

ADVISING & MENTORING

PhD Students

Student	Year	Description
Minbiao Han	2019 -	Passed qualification exam
Jibang Wu	2019 -	Passed qualification exam
Fan Yao	2020 -	Passed qualification exam, co-advised with Hongning Wang
Yang Yu	2021 -	From UVA Econ department,

Undergrads and Masters

Student	Year	Next Stop	Description
Harry Li	2022		UVA SEAS undergrad research fellowship
Emily Feng	2021 - 2022	Amazon	
Chenghan Zhou	2020 - 2022	Princeton	Leading-author papers at EC'22 and UAI'22; honorable mention for 2020 CRA Outstanding Undergrad Researcher award
Shuze Liu	2020 - 2022	PhD at UVA	Leading-author paper at EC'21
Quinlan Dawkins	2020 - 2022	nod.ai	Leading-author papers at ICML'21, NeurIPS'21 and AAAI'22
Nijat Khanbabayev	2021 - 2022	Hudson River Trading	
Sijun Tan	2020 - 2021	UC Berkeley	Leading-author paper at NeurIPS'21

Visitors

Student	Year	From	Description
James Nachbar	2020 summer	Yale	First-author paper at DAI'21

TEACHING

Lecturer

- UChicago CMSC 25300/35300: Math Foundations of Machine Learning (undergrad) Spring'23
- UChicago CMSC 27200: Theory of Algorithms (undergrad) Winter'23
- UChicago CMSC 35401: The Interplay of Learning and Game Theory (grad) Autumn'22
- UVA CS 4501: Introduction to Algorithmic Economics (undergrad) Spring'22
- UVA CS 4710: Artificial Intelligence (undergrad) Fall'21
- UVA CS 6501: Topics in Learning and Game Theory (grad) Spring'21
- UVA CS 6161: Design and Analysis of Algorithms (grad) Fall'20
- UVA CS 4710: Artificial Intelligence (undergrad) Spring'20
- UVA CS 6501: Topics in Learning and Game Theory (grad) Fall'19
- Harvard CS 182: Artificial Intelligence (undergrad) Fall'18

PROFESSIONAL SERVICE

Tutorials

- The Economics of Data and Learning (with James Zou and Shuran Zheng), AAAI 2023.
- Algorithmic Bayesian Persuasion (invited), the 17th Conference on Web and Internet Economics (WINE), Dec 2021.
- A Primer on the Interplay of Game Theory and Machine Learning, Summer School on Game Theory and Social Choice, CityU Hongkong, July 2021.
- Tutorial on Information, Persuasion and Decision Making at EC 2018.

Chairperson

- Co-organizer of IDEAL Special Quarter on Data Economics, Autumn 2022.
- Co-chair for the 13th Conference on Decision and Game Theory for Security (GameSec 2022)
- Workshop on Learning with Strategic Agents (LSA), with AAMAS 2022.
- Workshop on Strategic Reasoning for Societal Challenges (SRSC) with AAMAS 2019.
- Workshop on AI for Imperfect Information Games with AAAI 2018.
- Workshop on Adversarial Reasoning in Multi-Agent Systems with AAMAS 2017.

Senior Program Committee: AAAI (2019 – 2023), IJCAI (2021 – 2023)

Program Committee: EC(2019 – 2022), ICML (2022), NeurIPS (2021 – 2022), ICLR (2022), ALT (2022), AAAI (2018– 2019), IJCAI (2015 – 2020), WINE (2021), AAMAS (2019 – 2021), GameSec (2017 – 2019)

Journal Reviewing Activities: American Economic Review, Management Science, Mathematics of Operation Research, Games and Economic Behavior (GEB), Artificial Intelligence (AIJ), Journal of Artificial Intelligence Research (JAIR), Autonomous Agents and Multi-Agent Systems (JAAMAS), etc.

INVITED TALKS

[22]. *Towards an Efficient Market for Information.*, invited talk at Polytech Milan, Feb 2023

[21]. *Markov Persuasion Process and its Reinforcement Learning*, invited lecture at Harvard, Sep 2022

[20]. *Markov Persuasion Process and its Reinforcement Learning*, invited talk at the Institute of Automation, Chinese Academy of Sciences, Sep 2022

[19]. *Markov Persuasion Process and its Reinforcement Learning*, invited talk at the Simons Workshop on Quantifying Uncertainty, Berkeley, Sep 2022

[18]. *Optimal Pricing of Information*, invited talk at the Workshop on Data Value at UChicago, June 2022

[17]. *Algorithmic Information Design*, invited tutorial at the 17th Conference on Web and Internet Economics (WINE), Dec 2021

[16]. *Algorithmic Information Design: Computability, Learnability and Applicability to Societal Challenges*, the Center for Research on Computation and Society at Harvard, Nov 2021

[15]. *Algorithmic Information Design: Computability, Robustness and Learnability*, RPI Computer Science Colloquia & Seminars, Nov 2021

[14]. *Manipulating Learning Algorithms in Strategic Environments*, invited talk at Shanghai Finance and Economics University, August 2021

[13]. *Optimal Pricing of Information*, the Workshop on Strategic Communication and Learning, Stony Brook Center for Game Theory, July, 2021

[12]. *The Double-Edged Role of Information in Security Games*, Penn State University AI for Social Impact seminar series, Nov 2020

[11]. *The Double-Edged Role of Information in Security Games*, invited guest lecture at CMU for advanced course *AI Methods for Social Good*, Spring 2020

[10]. *Manipulating Learning Algorithms in Strategic Environments*, CMU Special Institute for Software Research Seminar, Feb 2020

[9]. *Manipulating Learning Algorithms in Strategic Environments*, UVA Bio-complexity Institute Seminar, Nov 2019

[8]. *Manipulating Learning Algorithms in Strategic Environments*, UVA Seminar on Artificial Intelligence and Machine Learning, Nov 2019

[7]. *Manipulating Learning Algorithms in Strategic Environments*, UVA Seminar on Human and Machine Intelligence, Oct 2019

[6]. *Algorithmic Persuasion with No Externalities*, Workshop on the Economics of Strategic Communication and Persuasion, Montreal, Nov 2018.

[5]. *Algorithmic Persuasion*, Harvard EconCS Seminar, Nov 2018.

[4]. *Information as A Double-Edged Sword in Strategic Interactions*, The Institute for Interdisciplinary Information Sciences (IIIS), Tsinghua University, June 2018.

[3]. *Strategic Coordination of Human Patrollers and UAVs with Signaling for Security Games*, Computational Sustainability Open Graduate Online Seminar, May 2018.

[2]. *The Mysteries of Security Games: Equilibrium Computation Becomes Combinatorial Algorithm Design*, Southern California Symposium on Network Economics and Game Theory (NEGT), Caltech, Jan 2018.

[1]. *Strategic Coordination of Human Patrollers and Mobile Sensors with Signaling for Security Games*, CMU CyLab, October 2017.